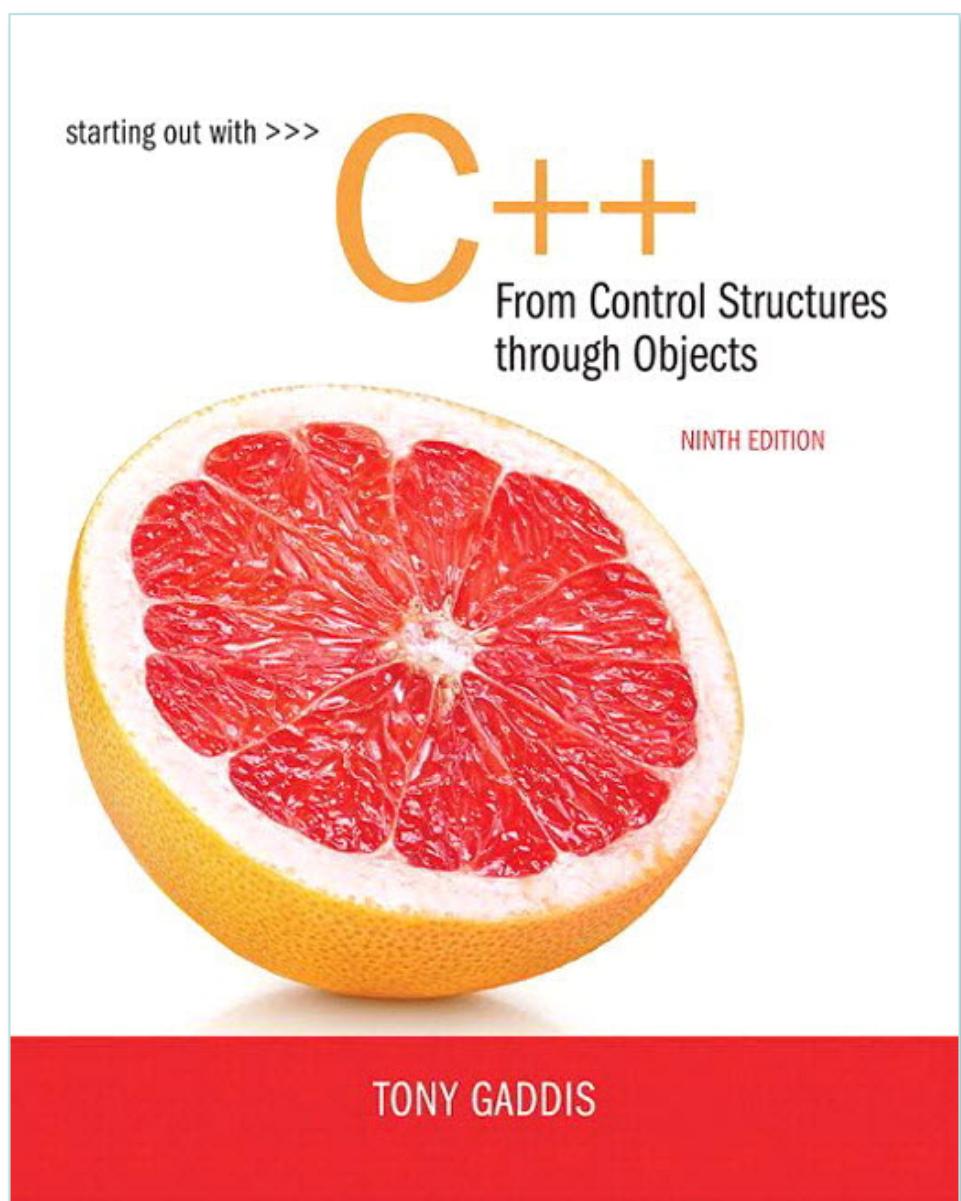
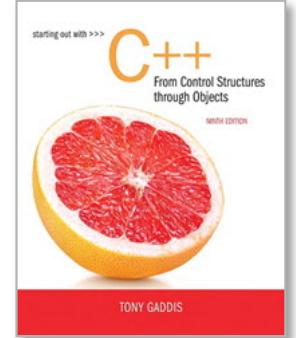


Chapter 17:

The Standard Template Library





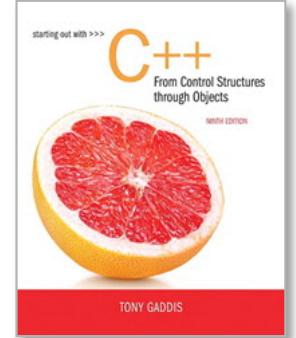
17.1

Introduction to the Standard Template Library

The Standard Template Library

- **The Standard Template Library (STL):** an extensive library of generic templates for classes and functions.
- **Categories of Templates:**
 - **Containers:** Class templates for objects that store and organize data
 - **Iterators:** Class templates for objects that behave like pointers, and are used to access the individual data elements in a container
 - **Algorithms:** Function templates that perform various operations on elements of containers





17.2

STL Container and Iterator Fundamentals

Containers

● Sequence Containers

- Stores data sequentially in memory, in a fashion similar to an array

● Associative Containers

- Stores data in a nonsequential way that makes it faster to locate elements



Containers

Table 17-1 Sequence Containers

Container Class	Description
array	A fixed-size container that is similar to an array
deque	A double-ended queue. Like a <code>vector</code> , but designed so that values can be quickly added to or removed from the front and back. (This container will be discussed in Chapter 19.)
forward_list	A singly linked list of data elements. Values may be inserted to or removed from any position. (This container will be discussed in Chapter 18.)
list	A doubly linked list of data elements. Values may be inserted to or removed from any position. (This container will be discussed in Chapter 18.)
vector	A container that works like an expandable array. Values may be added to or removed from a <code>vector</code> . The <code>vector</code> automatically adjusts its size to accommodate the number of elements it contains.



Containers

Table 17-2 Associative Containers

Container Class	Description
<code>set</code>	Stores a set of unique values that are sorted. No duplicates are allowed.
<code>multiset</code>	Stores a set of unique values that are sorted. Duplicates are allowed.
<code>map</code>	Maps a set of keys to data elements. Only one key per data element is allowed. Duplicates are not allowed. The elements are sorted in order of their keys.
<code>multimap</code>	Maps a set of keys to data elements. Many keys per data element are allowed. Duplicates are allowed. The elements are sorted in order of their keys.
<code>unordered_set</code>	Like a <code>set</code> , except that the elements are not sorted
<code>unordered_multiset</code>	Like a <code>multiset</code> , except that the elements are not sorted
<code>unordered_map</code>	Like a <code>map</code> , except that the elements are not sorted
<code>unordered_multimap</code>	Like a <code>multimap</code> , except that the elements are not sorted



Container Adapters

Table 17-3 Container Adapter Classes

Container Adapter Class	Description
stack	An adapter class that stores elements in a deque (by default). A stack is a last-in, first-out (LIFO) container. When you retrieve an element from a stack, the stack always gives you the last element that was inserted. (This class will be discussed in Chapter 19.)
queue	An adapter class that stores elements in a deque (by default). A queue is a first-in, first-out (FIFO) container. When you retrieve an element from a stack, the stack always gives you the first, or earliest, element that was inserted. (This class will be discussed in Chapter 19.)
priority_queue	An adapter class that stores elements in a vector (by default). A data structure in which the element that you retrieve is always the element with the greatest value. (This class will be discussed in Chapter 19.)



STL Header Files

Table 17-4 Header Files

Header File	Classes
<array>	array
<deque>	deque
<forward_list>	forward_list
<list>	list
<map>	map, multimap
<queue>	queue, priority_queue
<set>	set, multiset
<stack>	stack
<unordered_map>	unordered_map, unordered_multimap
<unordered_set>	unordered_set, unordered_multiset
<vector>	vector



The array Class Template

- An array object works very much like a regular array
- A fixed-size container that holds elements of the same data type.
- array objects have a `size()` member function that returns the number of elements contained in the object.

The array Class Template

- The array class is declared in the `<array>` header file.
- When defining an array object, you specify the data type of its elements, and the number of elements.
- Examples:

```
array<int, 5> numbers;
```

```
array<string, 4> names;
```

The array Class Template

- Initializing an array object:

```
array<int, 5> numbers = {1, 2, 3, 4, 5};
```

```
array<string, 4> names = {"Jamie", "Ashley", "Doug",
                           "Claire"};
```

The array Class Template

- The `array` class overloads the `[]` operator.
- You can use the `[]` operator to access elements using a subscript, just as you would with a regular array.
- The `[]` operator does not perform bounds checking. Be careful not to use a subscript that is out of bounds.

Program 17-1

```
1 #include <iostream>
2 #include <string>
3 #include <array>
4 using namespace std;
5
6 int main()
7 {
8     const int SIZE = 4;
9
10    // Store some names in an array object.
11    array<string, SIZE> names = {"Jamie", "Ashley", "Doug", "Claire"};
12
13    // Display the names.
14    cout << "Here are the names:\n";
15    for (int index = 0; index < names.size(); index++)
16        cout << names[index] << endl;
17
18    return 0;
19 }
```

Program Output

Here are the names:

Jamie
Ashley
Doug
Claire



Iterators

- Objects that work like pointers
- Used to access data in STL containers
- Five categories of iterators:

Table 17-6 Categories of Iterators

Iterator Category	Description
Forward	Can only move forward in a container (uses the <code>++</code> operator).
Bidirectional	Can move forward or backward in a container (uses the <code>++</code> and <code>--</code> operators).
Random access	Can move forward and backward, and can jump to a specific data element in a container.
Input	Can be used with an input stream to read data from an input device or a file.
Output	Can be used with an output stream to write data to an output device or a file.

Similarities between Pointers and Iterators

	Pointers	Iterators
Use the * and -> operators to dereference	Yes	Yes
Use the = operator to assign to an element	Yes	Yes
Use the == and != operators to compare	Yes	Yes
Use the ++ operator to increment	Yes	Yes
Use the -- operator to decrement	Yes	Yes (bidirectional and random-access iterators)
Use the + operator to move forward a specific number of elements	Yes	Yes
Use the - operator to move backward a specific number of elements	Yes	Yes Yes (bidirectional and random-access iterators)



Iterators

- To define an iterator, you must know what type of container you will be using it with.
- The general format of an iterator definition:

containerType::iterator iteratorName;

Where *containerType* is the STL container type, and *iteratorName* is the name of the iterator variable that you are defining.



Iterators

- For example, suppose we have defined an array object, as follows:

```
array<string, 3> names = {"Sarah", "William", "Alfredo"};
```

We can define an iterator that is compatible with the array object as follows:

```
array<string, 3>::iterator it;
```

This defines an iterator named `it`. The iterator can be used with an `array<string, 3>` object.



Iterators

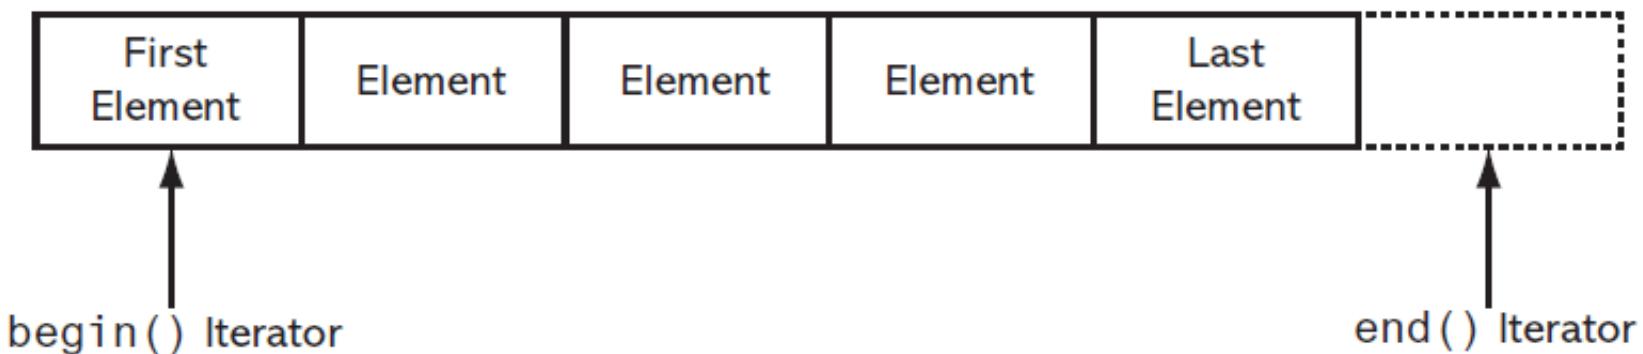
- ➊ All of the STL containers have a `begin()` member function that returns an iterator pointing to the container's first element.

```
// Define an array object.  
array<string, 3> names = {"Sarah", "William", "Alfredo"};  
  
// Define an iterator for the array object.  
array<string, 3>::iterator it;  
  
// Make the iterator point to the array object's first element.  
it = names.begin();  
  
// Display the element that the iterator points to.  
cout << *it << endl;
```



Iterators

- All of the STL containers have a `end()` member function that returns an iterator pointing to the position *after* the container's last element.



Iterators

- ➊ You typically use the `end()` member function to know when you have reached the end of a container.

```
// Define an array object.  
array<string, 3> names = {"Sarah", "William", "Alfredo"};  
  
// Define an iterator for the array object.  
array<string, 3>::iterator it;  
  
// Make the iterator point to the array object's first element.  
it = names.begin();  
  
// Display the array object's contents.  
while (it != names.end())  
{  
    cout << *it << endl;  
    it++;  
}
```



Program 17-2

```
1 #include <iostream>
2 #include <string>
3 #include <array>
4 using namespace std;
5
6 int main()
7 {
8     const int SIZE = 3;
9
10    // Store some names in an array object.
11    array<string, SIZE> names = {"Sarah", "William", "Alfredo"};
12
13    // Create an iterator for the array object.
14    array<string, SIZE>::iterator it;
15
16    // Display the names.
17    cout << "Here are the names:\n";
18    for (it = names.begin(); it != names.end(); it++)
19        cout << *it << endl;
20
21    return 0;
22 }
```

Program Output

Here are the names:
Sarah
William
Alfredo



Iterators

- You can use the `auto` keyword to simplify the definition of an iterator.
- Example:

```
array<string, 3> names = {"Sarah", "William", "Alfredo"};  
auto it = names.begin();
```

Program 17-3

```
1 #include <iostream>
2 #include <string>
3 #include <array>
4 using namespace std;
5
6 int main()
7 {
8     const int SIZE = 4;
9
10    // Store some names in an array object.
11    array<string, SIZE> names = {"Jamie", "Ashley", "Doug", "Claire"};
12
13    // Display the names.
14    cout << "Here are the names:\n";
15    for (auto it = names.begin(); it != names.end(); it++)
16        cout << *it << endl;
17
18    return 0;
19 }
```



Mutable Iterators

- ➊ An iterator of the `iterator` type gives you read/write access to the element to which the iterator points.
- ➋ This is commonly known as a mutable iterator.

```
// Define an array object.  
array<int, 5> numbers = {1, 2, 3, 4, 5};  
  
// Define an iterator for the array object.  
array<int, 5>::iterator it;  
  
// Make the iterator point to the array object's first element.  
it = numbers.begin();  
  
// Use the iterator to change the element.  
*it = 99;
```

Constant Iterators

- An iterator of the `const_iterator` type provides read-only access to the element to which the iterator points.
- The STL containers provide a `cbegin()` member function and a `cend()` member function.
 - The `cbegin()` member function returns a `const_iterator` pointing to the first element in a container.
 - The `cend()` member function returns a `const_iterator` pointing to the end of the container.
 - When working with `const_iterators`, simply use the container class's `cbegin()` and `cend()` member functions instead of the `begin()` and `end()` member functions.



Reverse Iterators

- A *reverse iterator* works in reverse, allowing you to iterate backward over the elements in a container.
- With a reverse iterator, the last element in a container is considered the first element, and the first element is considered the last element.
- The `++` operator moves a reverse iterator backward, and the `--` operator moves a reverse iterator forward.

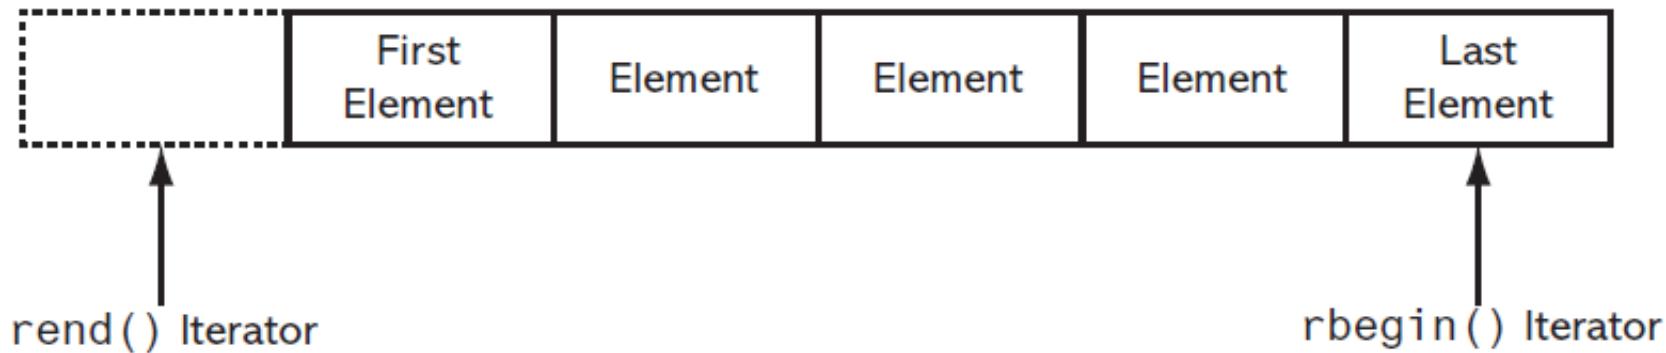


Reverse Iterators

- The following STL containers support reverse iterators:
 - array
 - deque
 - list
 - map
 - multimap
 - multiset
 - set
 - vector
- All of these classes provide an `rbegin()` member function and an `rend()` member function.

Reverse Iterators

- ➊ The `rbegin()` member function returns a reverse iterator pointing to the last element in a container.
- ➋ The `rend()` member function returns an iterator pointing to the position *before* the first element.

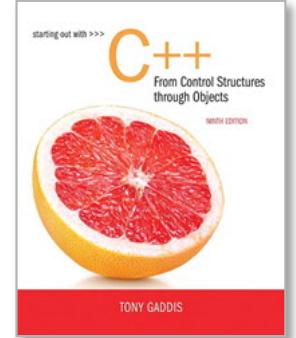


Reverse Iterators

- To create a reverse iterator, define it as `reverse_iterator`

```
// Define an array object.  
array<int, 5> numbers = {1, 2, 3, 4, 5};  
  
// Define a reverse iterator for the array object.  
array<int, 5>::reverse_iterator it;  
  
// Display the elements in reverse order.  
for (it = numbers.rbegin(); it != numbers.rend(); it++)  
    cout << *it << endl;
```





17.3

The vector Class

The vector Class

- A vector is a sequence container that works like an array, but is dynamic in size.
- Overloaded [] operator provides access to existing elements
- The vector class is declared in the <vector> header file.



vector Class Constructors

Default Constructor

`vector<dataType> name;`

Creates an empty vector.

Fill Constructor

`vector<dataType> name(size);`

Creates a vector of *size* elements. If the elements are objects, they are initialized via their default constructor. Otherwise, initialized with 0.

Fill Constructor

`vector<dataType> name(size, value);`

Creates a vector of *size* elements, each initialized with *value*.



vector Class Constructors

Range Constructor

`vector<dataType> name(iterator1, iterator2);`

Creates a vector that is initialized with a range of values from another container. *iterator1* marks the beginning of the range and *iterator2* marks the end.

Copy Constructor

`vector<dataType> name(vector2);`

Creates a vector that is a copy of *vector2*.



Program 17-4

```
1 #include <iostream>
2 #include <vector>
3 using namespace std;
4
5 int main()
6 {
7     const int SIZE = 10;
8
9     // Define a vector to hold 10 int values.
10    vector<int> numbers(SIZE);
11
12    // Store the values 0 through 9 in the vector.
13    for (int index = 0; index < numbers.size(); index++)
14        numbers[index] = index; ← Subscript notation
15
16    // Display the vector elements.
17    for (auto element : numbers) ← Range-based for loop
18        cout << element << " ";
19    cout << endl;
20
21    return 0;
22 }
```

Program Output

```
0 1 2 3 4 5 6 7 8 9
```



Initializing a vector

- In C++ 11 and later, you can initialize a vector object:

```
vector<int> numbers = {1, 2, 3, 4, 5};
```

or

```
vector<int> numbers {1, 2, 3, 4, 5};
```



Adding New Elements to a vector

- The `push_back` member function adds a new element to the end of a vector:

```
vector<int> numbers;  
numbers.push_back(10);  
numbers.push_back(20);  
numbers.push_back(30);
```

Accessing Elements with the at() Member Function

- >You can use the at() member function to retrieve a vector element by its index with bounds checking:

```
vector<string> names = {"Joe", "Karen", "Lisa"};
cout << names.at(0) << endl;
cout << names.at(1) << endl;
cout << names.at(2) << endl;
cout << names.at(3) << endl; // Throws an exception
```

Throws an out_of_bounds exception
when given an invalid index



Using an Iterator With a vector

- vectors have begin() and end() member functions that return iterators pointing to the beginning and end of the container:

```
// Create a vector containing names.  
vector<string> names = {"Joe", "Karen", "Lisa", "Jackie"};  
  
// Create an iterator.  
vector<string>::iterator it; // Defines an iterator that is compatible  
// Use the iterator to display each element in the vector.  
for (it = names.begin(); it != names.end(); it++)  
{  
    cout << *it << endl; // Displays the item that the iterator points to  
}
```



Using an Iterator With a vector

- The `begin()` and `end()` member functions return a random-access iterator of the `iterator` type
- The `cbegin()` and `cend()` member functions return a random-access iterator of the `const_iterator` type
- The `rbegin()` and `rend()` member functions return a reverse iterator of the `reverse_iterator` type
- The `crbegin()` and `crend()` member functions return a reverse iterator of the `const_reverse_iterator` type



Inserting Elements with the `insert()` Member Function

- You can use the `insert()` member function, along with an iterator, to insert an element at a specific position.
- General format:

```
vectorName.insert(it, value);
```

Iterator pointing to an element in the vector

Value to insert before the element that *it* points to



Program 17-5

```
1 #include <iostream>
2 #include <vector>
3 using namespace std;
4
5 int main()
6 {
7     // Define a vector with 5 int values.
8     vector<int> numbers = {1, 2, 3, 4, 5};
9
10    // Define an iterator pointing to the second element.
11    auto it = numbers.begin() + 1;
12
13    // Insert a new element with the value 99.
14    numbers.insert(it, 99);
15
16    // Display the vector elements.
17    for (auto element : numbers)
18        cout << element << " ";
19    cout << endl;
20
21    return 0;
22 }
```

Program Output

1 99 2 3 4 5



Overloaded Versions of the `insert()` Member Function

<code>insert(it, value)</code>	Inserts <i>value</i> just before the element pointed to by <i>it</i> . The function returns an iterator pointing to the newly inserted element.
<code>insert(it, n, value)</code>	Inserts <i>n</i> elements just before the element pointed to by <i>it</i> . Each of the new elements will be initialized with <i>value</i> . The function returns an iterator pointing to the first element of the newly inserted elements.
<code>insert(iterator1, iterator2, iterator3)</code>	Inserts a range of new elements. <i>iterator1</i> points to an existing element in the container. The range of new elements will be inserted before the element pointed to by <i>iterator1</i> . <i>iterator2</i> and <i>iterator3</i> mark the beginning and end of a range of values that will be inserted. (The element pointed to by <i>iterator3</i> will not be included in the range.) The function returns an iterator pointing to the first element of the newly inserted range.



Storing Objects Of Your Own Classes in a vector

- STL containers are especially useful for storing objects of your own classes.
- Consider this Product class:

```
1 #ifndef PRODUCT_H
2 #define PRODUCT_H
3 #include <string>
4 using namespace std;
5
6 class Product
7 {
8 private:
9     string name;
10    int units;
11 public:
12     Product(string n, int u)
13     { name = n;
14         units = u; }
15
16     void setName(string n)
17     { name = n; }
18
19     void setUnits(int u)
20     { units = u; }
21
22     string getName() const
23     { return name; }
24
25     int getUnits() const
26     { return units; }
27 };
28 #endif
```



Program 17-7

```
1 #include <iostream>
2 #include <vector>
3 #include "Product.h"
4 using namespace std;
5
6 int main()
7 {
8     // Create a vector of Product objects.
9     vector<Product> products =
10    {
11        Product("T-Shirt", 20),
12        Product("Calendar", 25),
13        Product("Coffee Mug", 30)
14    };
15
16    // Display the vector elements.
17    for (auto element : products)
18    {
19        cout << "Product: " << element.getName() << endl
20            << "Units: " << element.getUnits() << endl;
21    }
22
23    return 0;
24 }
```

Program Output

```
Product: T-Shirt
Units: 20
Product: Calendar
Units: 25
Product: Coffee Mug
Units: 30
```

This program initializes a vector with three Product objects.

A range-based for loop iterates over the vector.



Program 17-8

```
1 #include <iostream>
2 #include <string>
3 #include <vector>
4 #include "Product.h"
5 using namespace std;
6
7 int main()
8 {
9     // Create Product objects.
10    Product prod1("T-Shirt", 20);
11    Product prod2("Calendar", 25);
12    Product prod3("Coffee Mug", 30);
13
14    // Create a vector to hold the Products
15    vector<Product> products;
16
17    // Add the products to the vector.
18    products.push_back(prod1);
19    products.push_back(prod2);
20    products.push_back(prod3);
21
22    // Use an iterator to display the vector contents.
23    for (auto it = products.begin(); it != products.end(); it++)
24    {
25        cout << "Product: " << it->getName() << endl
26           << "Units: " << it->getUnits() << endl;
27    }
28
29    return 0;
30 }
```

This program uses the `push_back` member function to store three `Product` objects in a vector.

A for loop uses an iterator to step through the vector.

Program Output

```
Product: T-Shirt
Units: 20
Product: Calendar
Units: 25
Product: Coffee Mug
Units: 30
```



Inserting Container Elements With Emplacement

- Member functions such as `insert()` and `push_back()` can cause temporary objects to be created in memory while the insertion is taking place.
- This is not a problem in programs that make only a few insertions.
- However, these functions can be inefficient for making a lot of insertions.



Inserting Container Elements With Emplacement

- C++11 introduced a new family of member functions that use a technique known as *emplacement* to insert new elements.
- Emplacement avoids the creation of temporary objects in memory while a new object is being inserted into a container.
- The emplacement functions are more efficient than functions such as `insert()` and `push_back()`

Inserting Container Elements With Emplacement

- The vector class provides two member functions that use emplacement:
 - `emplace()` - emplaces an element at a specific location
 - `emplace_back()`- emplaces an element at the end of the vector
- With these member functions, it is not necessary to instantiate, ahead of time, the object you are going to insert.
- Instead, you pass to the emplacement function any arguments that you would normally pass to the constructor of the object you are inserting.
- The emplacement function handles the construction of the object, forwarding the arguments to its constructor.



Program 17-9

```
1 #include <iostream>
2 #include <vector>
3 #include "Product.h"
4 using namespace std;
5
6 int main()
7 {
8     // Create a vector to hold Products.
9     vector<Product> products; ← Define a vector to hold Product objects
10
11    // Add Products to the vector.
12    products.emplace_back("T-Shirt", 20);
13    products.emplace_back("Calendar", 25);
14    products.emplace_back("Coffee Mug", 30); ← Emplace three Product objects at the end of the vector
15
16    // Use an iterator to display the vector contents.
17    for (auto it = products.begin(); it != products.end(); it++)
18    {
19        cout << "Product: " << it->getName() << endl
20            << "Units: " << it->getUnits() << endl;
21    }
22
23    return 0;
24 }
```

Define a vector to hold Product objects

Emplace three Product objects at the end of the vector

A for loop uses an iterator to step through the vector.

Program Output

```
Product: T-Shirt
Units: 20
Product: Calendar
Units: 25
Product: Coffee Mug
Units: 30
```



Program 17-10

```
1 #include <iostream>
2 #include <vector>
3 #include "Product.h"
4 using namespace std;
5
6 int main()
7 {
8     // Create a vector to hold Products.
9     vector<Product> products =
10    {
11        Product("T-Shirt", 20),
12        Product("Coffee Mug", 30)
13    };
14
15    // Get an iterator to the 2nd element.
16    auto it = products.begin() + 1;
17
18    // Insert another Product into the vector.
19    products.emplace(it, "Calendar", 25);
20
21    // Display the vector contents.
22    for (auto element : products)
23    {
24        cout << "Product: " << element.getName() << endl
25            << "Units: " << element.getUnits() << endl;
26    }
27
28    return 0;
29 }
```

Initializes a vector with two Product objects

Gets an iterator pointing to the 2nd element

Emplaces a new Product object before the one pointed to by the iterator

Program Output

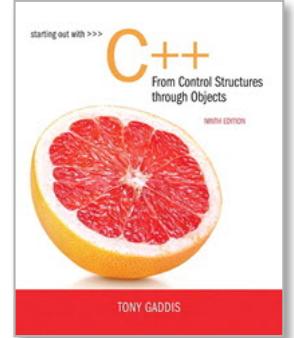
```
Product: T-Shirt
Units: 20
Product: Calendar
Units: 25
Product: Coffee Mug
Units: 30
```



The vector Class

- The vector class has many useful member functions.
- See Table 17-8 in your textbook.





17.4

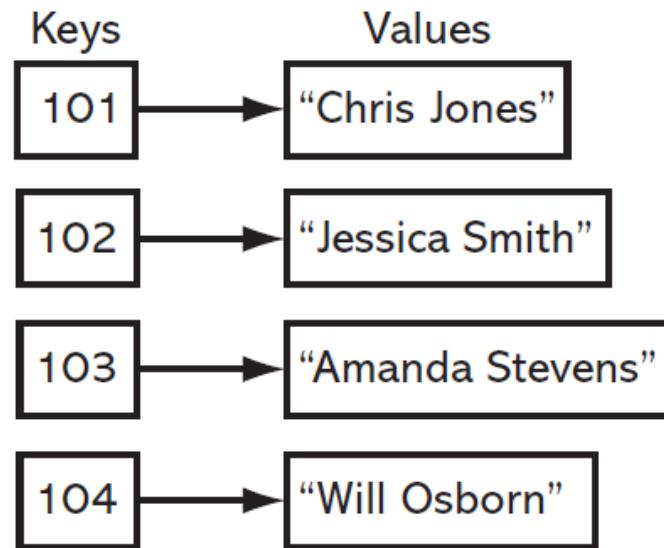
The `map`, `multimap`, and `unordered_map` Classes

Maps – General Concepts

- A *map* is an associative container.
- Each element that is stored in a map has two parts: a *key* and a *value*.
- To retrieve a specific value from a map, you use the key that is associated with that value.
- This is similar to the process of looking up a word in the dictionary, where the words are keys and the definitions are values.

Maps

- Example: a map in which employee IDs are the keys and employee names are the values.
- You use an employee's ID to look up that employee's name.



The map Class

- You can use the STL map class to store key-value pairs.
- The keys that are stored in a map container are unique – no duplicates.
- The map class is declared in the `<map>` header file.



map Class Constructors

Default Constructor `map<keyDataType, valueDataType> name;`
Creates an empty map.

Range Constructor `map<keyDataType, valueDataType>
name(iterator1, iterator2);`
Creates a map that is initialized with a range of values
from another map. *iterator1* marks the beginning of the
range and *iterator2* marks the end.

Copy Constructor `map<keyDataType, valueDataType> name(map2);`
Creates a map that is a copy of *map2*.



The map Class

- Example: defining a map container to hold employee ID numbers (as ints) and their corresponding employee names (as strings):

```
map<int, string> employees;
```

Key data type

Value data type

Initializing a map

```
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};
```

- In the first element, the key is 101 and the value is "Chris Jones".
- In the second element, the key is 102 and the value is "Jessica Smith".
- In the third element, the key is 103 and the value is "Amanda Stevens".
- In the fourth element, the key is 104 and the value is "Will Osborn".

The Overloaded [] Operator

- You can use the [] operator to add new elements to a map.
- General format:

mapName[key] = value;

- This adds the key-value pair to the map.
- If the key already exists in the map, its associated value will be changed to *value*.



The Overloaded [] Operator

```
map<int, string> employees;
employees[110] = "Beth Young";
employees[111] = "Jake Brown";
employees[112] = "Emily Davis";
```

- After this code executes, the employees map will contain the following elements:
 - Key = 110, Value = "Beth Young"
 - Key = 111, Value = "Jake Brown"
 - Key = 112, Value = "Emily Davis"



The pair Type

- Internally, the elements of a map are stored as instances of the pair type.
- pair is a struct that has two member variables: first and second.
- The element's key is stored in first, and the element's value is stored in second.
- The pair struct is declared in the <utility> header file. When you #include the <map> header file, <utility> is automatically included as well.



Inserting Elements with the `insert()` Member Function

- The `map` class provides an `insert()` member function that adds a `pair` object as an element to the `map`.
- You can use the STL function template `make_pair` to construct a `pair` object.
- The `make_pair` function template is declared in the `<utility>` header file.

Inserting Elements with the `insert()` Member Function

```
map<int, string> employees;
employees.insert(make_pair(110, "Beth Young"));
employees.insert(make_pair(111, "Jake Brown"));
employees.insert(make_pair(112, "Emily Davis"));
```

- After this code executes, the `employees` map will contain the following elements:
 - Key = 110, Value = "Beth Young"
 - Key = 111, Value = "Jake Brown"
 - Key = 112, Value = "Emily Davis"

Note: If the element that you are inserting with the `insert()` member function has the same key as an existing element, the function will *not* insert the new element.

Inserting Elements with the `emplace()` Member Function

- The map class also provides an `emplace()` member function that adds an element to the map.

```
map<int, string> employees;
employees.emplace(110, "Beth Young");
employees.emplace(111, "Jake Brown");
employees.emplace(112, "Emily Davis");
```

- After this code executes, the `employees` map will contain the following elements:
 - Key = 110, Value = "Beth Young"
 - Key = 111, Value = "Jake Brown"
 - Key = 112, Value = "Emily Davis"

Note: If the element that you are inserting with the `emplace()` member function has the same key as an existing element, the function will *not* insert the new element.

Retrieving Elements with the at() Member Function

- You can use the at() member function to retrieve a map element by its key:

```
// Create a map containing employee IDs and names.  
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};  
  
// Retrieve a value from the map.  
cout << employees.at(103) << endl;
```

Displays "Amanda Stevens"



Retrieving Elements with the at() Member Function

- To prevent the at() member function from throwing an exception (if the specified key does not exist), use the count member function to determine whether it exists:

```
// Create a map containing employee IDs and names.  
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};  
  
// Retrieve a value from the map.  
if (employees.count(103)) ←———— The count() member function  
    cout << employees.at(103) << endl;  
else  
    cout << "Employee not found.\n";
```

Deleting Elements

- 💡 You can use the `erase()` member function to retrieve a map element by its key:

```
// Create a map containing employee IDs and names.  
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};  
  
// Delete the employee with ID 102.  
employees.erase(102);
```

Deletes Jessica Smith from the map



Stepping Through a map with the Range-Based for Loop

```
// Create a map containing employee IDs and names.  
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};  
  
// Display each element.  
for (pair<int, string> element : employees)  
{  
    cout << "ID: " << element.first << "\tName: "  
        << element.second << endl;  
}
```

Remember, each element is a pair.



Stepping Through a map with the Range-Based for Loop

```
// Create a map containing employee IDs and names.  
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};  
  
// Display each element.  
for (auto element : employees) ← auto simplifies this  
{  
    cout << "ID: " << element.first << "\tName: "  
        << element.second << endl;  
}
```

Using an Iterator With a map

- The `begin()` and `end()` member functions return a bidirectional iterator of the `iterator` type
- The `cbegin()` and `cend()` member functions return a bidirectional iterator of the `const_iterator` type
- The `rbegin()` and `rend()` member functions return a reverse bidirectional iterator of the `reverse_iterator` type
- The `crbegin()` and `crend()` member functions return a reverse bidirectional iterator of the `const_reverse_iterator` type

Using an Iterator With a map

- When an iterator points to a `map` element, it points to an instance of the `pair` type.
- The element has two member variables: `first` and `second`.
- The element's key is stored in `first`, and the element's value is stored in `second`.

Program 17-11

```
1 // This program demonstrates an iterator with a map.
2 #include <iostream>
3 #include <string>
4 #include <map>
5 using namespace std;
6
7 int main()
8 {
9     // Create a map containing employee IDs and names.
10    map<int, string> employees =
11        { {101,"Chris Jones"}, {102,"Jessica Smith"} ,
12          {103,"Amanda Stevens"},{104,"Will Osborn"} } ;
13
14    // Create an iterator.
15    map<int, string>::iterator iter;
16
17    // Use the iterator to display each element in the map.
18    for (iter = employees.begin(); iter != employees.end(); iter++)
19    {
20        cout << "ID: " << iter->first
21            << "\tName: " << iter->second << endl;
22    }
23
24    return 0;
25 }
```

Program Output

ID: 101 Name: Chris Jones
ID: 102 Name: Jessica Smith
ID: 103 Name: Amanda Stevens
ID: 104 Name: Will Osborn



Storing Objects Of Your Own Classes as *Values* in a map

- If you want to store an object as a value in a map, there is one requirement for that object's class:

It must have a default constructor.

- Consider the following Contact class...

```
1 #ifndef CONTACT_H
2 #define CONTACT_H
3 #include <string>
4 using namespace std;
5
6 class Contact
7 {
8 private:
9     string name;
10    string email;
11 public:
12     Contact()           ← Default constructor
13     {   name = "";      ←
14         email = "";   ← }
15
16     Contact(string n, string em)
17     {   name = n;
18         email = em;   ← }
19
20     void setName(string n)
21     {   name = n;   ← }
22
23     void setEmail(string em)
24     {   email = em;   ← }
25
26     string getName() const
27     {   return name;   ← }
28
29     string getEmail() const
30     {   return email;   ← }
31 };
32 #endif
```



Program 17-14

```
1 #include <iostream>
2 #include <string>
3 #include <map>
4 #include "Contact.h"
5 using namespace std;
6
7 int main()
8 {
9     string searchName;    // The name to search for
10
11    // Create some Contact objects
12    Contact contact1("Ashley Miller", "amiller@faber.edu");
13    Contact contact2("Jacob Brown", "jbrown@gotham.edu");
14    Contact contact3("Emily Ramirez", "eramirez@coolidge.edu");
15
16    // Create a map to hold the Contact objects.
17    map<string, Contact> contacts;
18
19    // Create an iterator for the map.
20    map<string, Contact>::iterator iter;
21
22    // Add the contact objects to the map.
23    contacts[contact1.getName()] = contact1;
24    contacts[contact2.getName()] = contact2;
25    contacts[contact3.getName()] = contact3;
```

In the map, the keys are the contact names, and the values are the Contact objects.



```
26
27     // Get the name to search for.
28     cout << "Enter a name: ";
29     getline(cin, searchName);
30
31     // Search for the name.
32     iter = contacts.find(searchName);
33
34     // Display the results.
35     if (iter != contacts.end())
36     {
37         cout << "Name: " << iter->second.getName() << endl;
38         cout << "Email: " << iter->second.getEmail() << endl;
39     }
40     else
41     {
42         cout << "Contact not found.\n";
43     }
44
45     return 0;
46 }
```

Program Output (with Example Input Shown in Bold)

```
Enter a name: Emily Ramirez 
Name: Emily Ramirez
Email: eramirez@coolidge.edu
```

Program Output (with Example Input Shown in Bold)

```
Enter a name: Billy Clark 
Contact not found.
```



Storing Objects Of Your Own Classes as Keys in a map

- If you want to store an object as a key in a map, there is one requirement for that object's class:

It must overload the < operator.

- Consider the following Customer class...

```
1 #ifndef CUSTOMER_H
2 #define CUSTOMER_H
3 #include<string>
4 using namespace std;
5
6 class Customer
7 {
8 private:
9     int custNumber;
10    string name;
11 public:
12     Customer(int cn, string n)
13     { custNumber = cn;
14         name = n; }
15
16     void setCustNumber(int cn)
17     { custNumber = cn; }
18
19     void setName(string n)
20     { name = n; }
21
22     int getCustNumber() const
23     { return custNumber; }
24
25     string getName() const
26     { return name; }
27
28     bool operator < (const Customer &right) const
29     { bool status = false;
30
31         if (custNumber < right.custNumber)
32             status = true;
33
34         return status; }
35     };
36 #endif
```



Program 17-17

```
1 #include <iostream>
2 #include <string>
3 #include <map>
4 #include "Customer.h"
5 using namespace std;
6
7 int main()
8 {
9     // Create some Customer objects.
10    Customer customer1(1001, "Sarah Scott");
11    Customer customer2(1002, "Austin Hill");
12    Customer customer3(1003, "Megan Cruz");
13
14    // Create a map to hold the seat assignments.
15    map<Customer, string> assignments;
16
17    // Use the map to store the seat assignments.
18    assignments[customer1] = "1A";
19    assignments[customer2] = "2B";
20    assignments[customer3] = "3C";
21
22    // Display all objects in the map.
23    for (auto element : assignments)
24    {
25        cout << element.first.getName() << "\t"
26            << element.second << endl;
27    }
28
29    return 0;
30 }
```

Program Output

Sarah Scott	1A
Austin Hill	2B
Megan Cruz	3C

This program assigns seats in a theater to customers. The map uses Customer objects as keys, and seat numbers as values.



The `unordered_map` Class

- The `unordered_map` class is similar to the `map` class, except in two regards:
 - The keys in an `unordered_map` are not sorted
 - The `unordered_map` class has better performance
- You should use the `unordered_map` class instead of the `map` class if:
 - You will be making a lot of searches on a large number of elements
 - You are not concerned with retrieving them in key order
- The `unordered_map` class is declared in the `<unordered_map>` header file



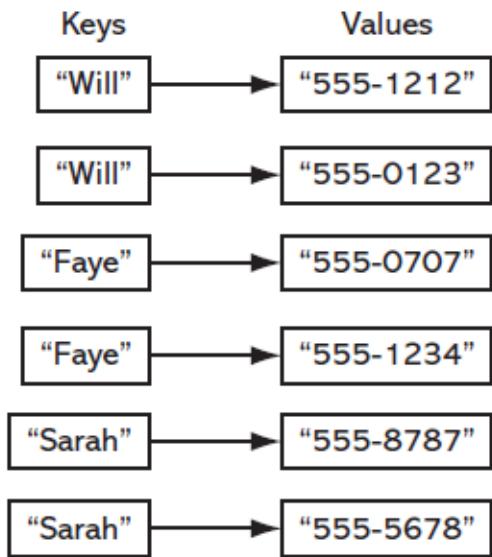
The multimap Class

- The multimap class is a map that allows duplicate keys
- The multimap class has most of the same member functions as the map class (see Table 17-11 in your textbook)
- The multimap class is declared in the `<map>` header file



The multimap Class

- Consider a phonebook application where the key is a person's name and the value is that person's phone number.
- A multimap container would allow each person to have multiple phone numbers



Program 17-19

```
1 #include <iostream>
2 #include <string>
3 #include <map>
4 using namespace std;
5
6 int main()
7 {
8     // Define a phonebook multimap.
9     multimap<string, string> phonebook =
10        { {"Will", "555-1212"}, {"Will", "555-0123"},
11         {"Faye", "555-0707"}, {"Faye", "555-1234"},
12         {"Sarah", "555-8787"}, {"Sarah", "555-5678"} };
13
14    // Display the elements in the multimap.
15    for (auto element : phonebook)
16    {
17        cout << element.first << "\t"
18            << element.second << endl;
19    }
20    return 0;
21 }
```

Program Output

Faye 555-0707
Faye 555-1234
Sarah 555-8787
Sarah 555-5678
Will 555-1212
Will 555-0123



Adding Elements to a multimap

- The multimap class does not overload the [] operator.
 - So, you cannot use an assignment statement to add a new element to a multimap.
- Instead, you will use either the emplace() or the insert() member functions.



Adding Elements to a multimap

```
multimap<string, string> phonebook;  
phonebook.emplace("Will", "555-1212");  
phonebook.emplace("Will", "555-0123");  
phonebook.emplace("Faye", "555-0707");  
phonebook.emplace("Faye", "555-1234");  
phonebook.emplace("Sarah", "555-8787");  
phonebook.emplace("Sarah", "555-5678");
```

Adding Elements to a multimap

```
multimap<string, string> phonebook;
phonebook.insert(make_pair("Will", "555-1212"));
phonebook.insert(make_pair("Will", "555-0123"));
phonebook.insert(make_pair("Faye", "555-0707"));
phonebook.insert(make_pair("Faye", "555-1234"));
phonebook.insert(make_pair("Sarah", "555-8787"));
phonebook.insert(make_pair("Sarah", "555-5678"));
```

Getting the Number of Elements With a Specified Key

- The `multimap` class's `count()` member function accepts a key as its argument, and returns the number of elements that match the specified key.

Program 17-20

```
1 #include <iostream>
2 #include <string>
3 #include <map>
4 using namespace std;
5
6 int main()
7 {
8     // Define a phonebook multimap.
9     multimap<string, string> phonebook =
10     { {"Will", "555-1212"}, {"Will", "555-0123"} ,
11      {"Faye", "555-0707"}, {"Faye", "555-1234"} ,
12      {"Sarah", "555-8787"}, {"Sarah", "555-5678"} };
13
14     // Display the number of elements that match "Faye".
15     cout << "Faye has " << phonebook.count("Faye") << " elements.\n";
16     return 0;
17 }
```

Program Output

Faye has 2 elements.



Retrieving Elements with a Specified Key

- The `multimap` class has a `find()` member function that searches for an element with a specified key.
- The `find()` function returns an iterator to the first element matching it.
- If the element is not found, the `find()` function returns an iterator to the end of the `multimap`.

Retrieving Elements with a Specified Key

- To retrieve all elements matching a specified key, use the `equal_range` member function.
- The `equal_range` member function returns a pair object.
 - The pair object's first member is an iterator pointing to the first element that matches the specified key.
 - The pair object's second member is an iterator pointing to the position *after* the last element that matches the specified key.



```
// Define a phonebook multimap.  
multimap<string, string> phonebook =  
{ {"Will", "555-1212"}, {"Will", "555-0123"},  
  {"Faye", "555-0707"}, {"Faye", "555-1234"},  
  {"Sarah", "555-8787"}, {"Sarah", "555-5678"} };  
  
// Define a pair variable to receive the object that  
// is returned from the equal_range member function.  
pair<multimap<string, string>::iterator,  
      multimap<string, string>::iterator> range;  
  
// Define an iterator for the multimap.  
multimap<string, string>::iterator iter;  
  
// Get the range of elements that match "Faye".  
range = phonebook.equal_range("Faye");  
  
// Display all of the elements that match "Faye".  
for (iter = range.first; iter != range.second; iter++)  
{  
    cout << iter->first << "\t" << iter->second << endl;  
}
```

Program Output

```
Faye      555-0707  
Faye      555-1234
```



Deleting Elements with a Specified Key

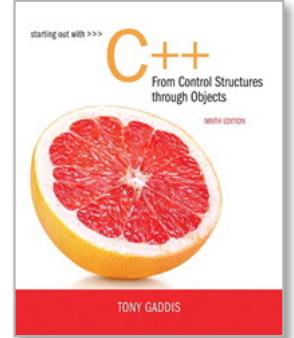
- To delete all elements matching a specified key, use the `erase()` member function.

```
// Define a phonebook multimap.  
multimap<string, string> phonebook =  
    { {"Will", "555-1212"}, {"Will", "555-0123"},  
     {"Faye", "555-0707"}, {"Faye", "555-1234"},  
     {"Sarah", "555-8787"}, {"Sarah", "555-5678"} };  
  
// Delete Will's phone numbers from the multimap.  
phonebook.erase("Will");
```

The unordered_multimap Class

- The `unordered_multimap` class is similar to the `multimap` class, except:
 - The keys in an `unordered_multimap` are not sorted
 - The `unordered_multimap` class has better performance
- You should use the `unordered_multimap` class instead of the `multimap` class if:
 - You will be making a lot of searches on a large number of elements
 - You are not concerned with retrieving them in key order
- The `unordered_multimap` class is declared in the `<unordered_multimap>` header file





17.5

The `set`, `multiset`, and `unordered_set` Classes

Sets

- A *set* is an associative container that is similar to a mathematical set.
- You can use the STL set class to create a set container.
- All the elements in a set must be unique. No two elements can have the same value.
- The elements in a set are automatically sorted in ascending order.
- The set class is declared in the `<set>` header file.



The set Class

- You can use the STL set class to create a set container.
- The keys that are stored in a map container are unique – no duplicates.
- The map class is declared in the `<map>` header file.



set Class Constructors

Default Constructor `set<dataType> name;`
Creates an empty set.

Range Constructor `set<dataType> name(iterator1, iterator2);`
Creates a set that is initialized with a range of values.
iterator1 marks the beginning of the range and
iterator2 marks the end.

Copy Constructor `set<dataType> name(set2);`
Creates a set that is a copy of *set2*.



The set Class

- Example: defining a set container to hold integers:

```
set<int> numbers;
```

- Example: defining and initializing a set container to hold integers:

```
set<int> numbers = {1, 2, 3, 4, 5};
```

The set Class

- A set cannot contain duplicate items.
- If the same value appears more than once in an initialization list, it will be added to the set only one time.
- For example, the following set will contain the values 1, 2, 3, 4, and 5:

```
set<int> numbers = {1, 1, 2, 2, 3, 4, 5, 5, 5};
```

Adding New Elements to a set

- The `insert()` member function adds a new element to a set:

```
set<int> numbers;  
numbers.insert(10);  
numbers.insert(20);  
numbers.insert(30);
```

Stepping Through a set With the Range-Based for Loop

```
// Create a set containing names.  
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};  
  
// Display each element.  
for (string element : names)  
{  
    cout << element << endl;  
}
```

Using an Iterator With a `set`

- The `begin()` and `end()` member functions return a bidirectional iterator of the `iterator` type
- The `cbegin()` and `cend()` member functions return a bidirectional iterator of the `const_iterator` type
- The `rbegin()` and `rend()` member functions return a reverse bidirectional iterator of the `reverse_iterator` type
- The `crbegin()` and `crend()` member functions return a reverse bidirectional iterator of the `const_reverse_iterator` type



Using an Iterator With a set

```
// Create a set containing names.  
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};  
  
// Create an iterator.  
set<string>::iterator iter;  
  
// Use the iterator to display each element in the set.  
for (iter = names.begin(); iter != names.end(); iter++)  
{  
    cout << *iter << endl;  
}
```

Determining Whether an Element Exists

- The set class's `count()` member function accepts a value as its argument, and returns 1 if that value exists in the set. The function returns 0 otherwise.

```
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};
if (names.count("Lisa"))
    cout << "Lisa was found in the set.\n";
else
    cout << "Lisa was not found.\n";
```

Retrieving an Element

- The set class has a `find()` member function that searches for an element with a specified value.
- The `find()` function returns an iterator to the element matching it.
- If the element is not found, the `find()` function returns an iterator to the end of the set.

Retrieving an Element

```
// Create a set containing names.  
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};  
  
// Create an iterator.  
set<string>::iterator iter;  
  
// Find "Karen".  
iter = names.find("Karen");  
  
// Display the result.  
if (iter != names.end())  
{  
    cout << *iter << " was found.\n";  
}  
else  
{  
    cout << "Karen was not found.\n";  
}
```



Storing Objects Of Your Own Classes in a set

- If you want to store an object in a set, there is one requirement for that object's class:

It must overload the < operator.

- Consider the following Customer class...

```
1 #ifndef CUSTOMER_H
2 #define CUSTOMER_H
3 #include<string>
4 using namespace std;
5
6 class Customer
7 {
8 private:
9     int custNumber;
10    string name;
11 public:
12     Customer(int cn, string n)
13     { custNumber = cn;
14         name = n; }
15
16     void setCustNumber(int cn)
17     { custNumber = cn; }
18
19     void setName(string n)
20     { name = n; }
21
22     int getCustNumber() const
23     { return custNumber; }
24
25     string getName() const
26     { return name; }
27
28     bool operator < (const Customer &right) const
29     { bool status = false;
30
31         if (custNumber < right.custNumber)
32             status = true;
33
34         return status; }
35     };
36 #endif
```



Program 17-22

```
1 #include <iostream>
2 #include <set>
3 #include "Customer.h"
4 using namespace std;
5
6 int main()
7 {
8     // Create a set of Customer objects.
9     set<Customer> customerset =
10    { Customer(1003, "Megan Cruz"),
11      Customer(1002, "Austin Hill"),
12      Customer(1001, "Sarah Scott")
13    };
14
15    // Try to insert a duplicate customer number.
16    customerset.emplace(1001, "Evan Smith");
17
18    // Display the set elements
19    cout << "List of customers:\n";
20    for (auto element : customerset)
21    {
22        cout << element.getCustNumber() << " "
23            << element.getName() << endl;
24    }
25
```

Continued...



```
26 // Search for customer number 1002.  
27 cout << "\nSearching for Customer Number 1002:\n";  
28 auto it = customerset.find(Customer(1002, ""));  
29  
30 if (it != customerset.end())  
31     cout << "Found: " << it->getName() << endl;  
32 else  
33     cout << "Not found.\n";  
34  
35 return 0;  
36 }
```

Program Output

List of customers:

1001 Sarah Scott
1002 Austin Hill
1003 Megan Cruz

Searching for Customer Number 1002:

Found: Austin Hill



The multiset Class

- The multiset class is a set that allows duplicate items.
- The multiset class has the same member functions as the set class (see Table 17-13 in your textbook).
- The multiset class is declared in the `<set>` header file.



The multiset Class

- In the set class, the count() member function returns either 0 or 1. In the multiset class, the count() member function can return values greater than 1.
- In the set class, the equal_range() member function returns a range with, at most, one element. In the multiset class, the equal_range() member function can return a range with multiple elements.



The `unordered_set` Class

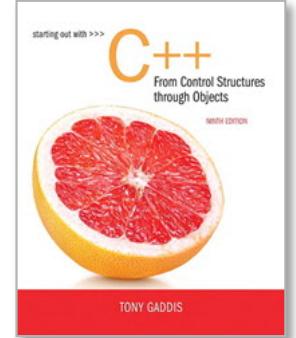
- The `unordered_set` class is similar to the `set` class, except in two regards:
 - The values in an `unordered_set` are not sorted
 - The `unordered_set` class has better performance
- You should use the `unordered_set` class instead of the `set` class if:
 - You will be making a lot of searches on a large number of elements
 - You are not concerned with retrieving them in ascending order
- The `unordered_set` class is declared in the `<unordered_set>` header file



The `unordered_multiset` Class

- The `unordered_multiset` class is similar to the `multiset` class, except in two regards:
 - The values in an `unordered_multiset` are not sorted
 - The `unordered_multiset` class has better performance
- You should use the `unordered_multiset` class instead of the `multiset` class if:
 - You will be making a lot of searches on a large number of elements
 - You are not concerned with retrieving them in ascending order
- The `unordered_multiset` class is declared in the `<unordered_set>` header file





17.6

Algorithms

STL Algorithms

- The STL provides a number of algorithms, implemented as function templates, in the `<algorithm>` header file.
- These functions perform various operations on ranges of elements.
- A range of elements is a sequence of elements denoted by two iterators:
 - The first iterator points to the first element in the range
 - The second iterator points to the end of the range (the element to which the second iterator points is not included in the range).



Categories of Algorithms in the STL

- Min/max algorithms
- Sorting algorithms
- Search algorithms
- Read-only sequence algorithms
- Copying and moving algorithms
- Swapping algorithms
- Replacement algorithms
- Removal algorithms
- Reversal algorithms
- Fill algorithms
- Rotation algorithms
- Shuffling algorithms
- Set algorithms
- Transformation algorithm
- Partition algorithms
- Merge algorithms
- Permutation algorithms
- Heap algorithms
- Lexicographical comparison algorithm



Sorting

- The sort function:

```
sort(iterator1, iterator2);
```

iterator1 and *iterator2* mark the beginning and end of a range of elements. The function sorts the range of elements in ascending order.



Searching

- The `binary_search` function:

```
binary_search(iterator1, iterator2, value);
```

iterator1 and *iterator2* mark the beginning and end of a range of elements that are sorted in ascending order. *value* is the value to search for. The function returns `true` if *value* is found in the range, or `false` otherwise.



Program 17-23

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4 using namespace std;
5
6 int main()
7 {
8     int searchValue;    // Value to search for
9
10    // Create a vector of unsorted integers.
11    vector<int> numbers = {10, 1, 9, 2, 8, 3, 7, 4, 6, 5};
12
13    // Sort the vector.
14    sort(numbers.begin(), numbers.end());
15
```

Continued...



```
16     // Display the vector.  
17     cout << "Here are the sorted values:\n";  
18     for (auto element : numbers)  
19         cout << element << " ";  
20     cout << endl;  
21  
22     // Get the value to search for.  
23     cout << "Enter a value to search for: ";  
24     cin >> searchValue;  
25  
26     // Search for the value.  
27     if (binary_search(numbers.begin(), numbers.end(), searchValue))  
28         cout << "That value is in the vector.\n";  
29     else  
30         cout << "That value is not in the vector.\n";  
31  
32     return 0;  
33 }
```

Program Output

Here are the sorted values:
1 2 3 4 5 6 7 8 9 10
Enter a value to search for: 8
That value is in the vector.

Program Output

Here are the sorted values:
1 2 3 4 5 6 7 8 9 10
Enter a value to search for: 99
That value is not in the vector.



Detecting Permutations

- If a range has N elements, there are $N!$ possible arrangements, or permutations, of those elements.
- For example, the range of integers 1, 2, 3 has six possible permutations:

1 , 2 , 3
1 , 3 , 2
2 , 1 , 3
2 , 3 , 1
3 , 1 , 2
3 , 2 , 1



Detecting Permutations

- The `is_permutation()` function determines whether one range of elements is a permutation of another range of elements.

`is_permutation(iterator1, iterator2, iterator3)`

- `iterator1` and `iterator2` mark the beginning and end of the first range of elements.
- `iterator3` marks the beginning of the second range of elements, assumed to have the same number of elements as the first range.
- The function returns true if the second range is a permutation of the first range, or false otherwise.
- See Program 17-25 in your textbook for an example.

Plugging Your Own Functions into an Algorithm

- Many of the function templates in the STL are designed to accept function pointers as arguments.
- This allows you to “plug” one of your own functions into the algorithm.
- For example:

```
for_each(iterator1, iterator2, function)
```

- *iterator1* and *iterator2* mark the beginning and end of a range of elements.
- *function* is the name of a function that accepts an element as its argument.
- The `for_each()` function iterates over the range of elements, passing each element as an argument to *function*.



Plugging Your Own Functions into an Algorithm

- For example, consider this function:

```
void doubleNumber(int &n)
{
    n = n * 2;
}
```

Plugging Your Own Functions into an Algorithm

And this code snippet:

```
vector<int> numbers = { 1, 2, 3, 4, 5 };

// Display the numbers before doubling.
for (auto element : numbers)
    cout << element << " ";
cout << endl;

// Double the value of each vector element.
for_each(numbers.begin(), numbers.end(), doubleNumber);

// Display the numbers before doubling.
for (auto element : numbers)
    cout << element << " ";
cout << endl;
```

This passes each element of the numbers vector to the doubleNumber function.



Plugging Your Own Functions into an Algorithm

- Another example:

`count_if(iterator1, iterator2, function)`

- *iterator1* and *iterator2* mark the beginning and end of a range of elements.
- *function* is the name of a function that accepts an element as its argument, and returns either true or false.
- The `count_if()` function iterates over the range of elements, passing each element as an argument to *function*.
- The `count_if` function returns the number of elements for which *function* returns true.



```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4 using namespace std;
5
6 // Function prototypes
7 bool isNegative(int);
8
9 int main()
10 {
11     // Create a vector of ints.
12     vector<int> numbers = { 0, 99, 120, -33, 10, 8, -1, 101 };
13
14     // Get the number of elements that are negative.
15     int negatives = count_if(numbers.begin(), numbers.end(), isNegative);
16
17     // Display the results.
18     cout << "There are " << negatives << " negative elements.\n";
19     return 0;
20 }
21
22 // isNegative function
23 bool isNegative(int n)
24 {
25     bool status = false;
26
27     if (n < 0)
28         status = true;
29
30     return status;
31 }
```

Program Output

There are 2 negative elements.

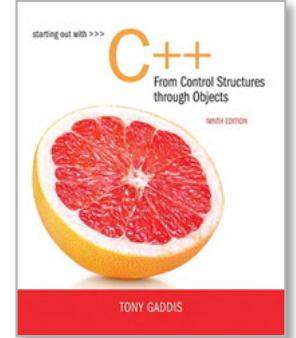


Algorithms for Set Operations

- The STL provides function templates for basic mathematical set operations.

STL Function Template	Description
<code>set_union</code>	Finds the union of two sets, which is a set that contains all the elements of both sets, excluding duplicates.
<code>set_intersection</code>	Finds the intersection of two sets, which is a set that contains only the elements that are found in both sets.
<code>set_difference</code>	Finds the difference of two sets, which is the set of elements that appear in one set, but not the other.
<code>set_symmetric_difference</code>	Finds the symmetric difference of two sets, which is the set of elements that appear in one set, but not both.
<code>set_includes</code>	Determines whether one set includes another.





17.7

Introduction to Function Objects and Lambda Expressions

Function Objects

- A function object is an object that acts like a function.
 - It can be called
 - It can accept arguments
 - It can return a value
- Function objects are also known as *functors*



Function Objects

- To create a function object, you write a class that overloads the () operator.

```
1 #ifndef SUM_H
2 #define SUM_H
3
4 class Sum
5 {
6 public:
7     int operator()(int a, int b)
8     { return a + b; } ← Returns an int
9 }
10#endif
```

Accepts two int arguments



Program 17-34

```
1 #include <iostream>
2 #include "Sum.h"
3 using namespace std;
4
5 int main()
6 {
7     // Local variables
8     int x = 10;
9     int y = 2;
10    int z = 0;
11
12    // Create a Sum object.
13    Sum sum;
14
15    // Call the sum function object.
16    z = sum(x, y);
17
18    // Display the result.
19    cout << z << endl;
20
21    return 0;
22 }
```

Program Output

12



Anonymous Function Objects

- Function objects can be called at the point of their creation, without being given a name. Consider this class:

```
1 #ifndef IS EVEN_H
2 #define IS EVEN_H
3
4 class IsEven
5 {
6 public:
7     bool operator()(int x)
8     { return x % 2 == 0; }
9 };
10#endif
```

Program 17-36

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4 #include "IsEven.h"
5 using namespace std;
6
7 int main()
8 {
9     // Create a vector of ints.
10    vector<int> v = { 1, 2, 3, 4, 5, 6, 7, 8 };
11
12    // Get the number of elements that even.
13    int evenNums = count_if(v.begin(), v.end(), IsEven());
14
15    // Display the results.
16    cout << "The vector contains " << evenNums << " even numbers.\n";
17    return 0;
18 }
```

An IsEven object is created here, but not given a name.
It is anonymous.



Program Output

The vector contains 4 even numbers.



Predicate Terminology

- A function or function object that returns a Boolean value is called a *predicate*.
- A predicate that takes only one argument is called a *unary predicate*.
- A predicate that takes two arguments is called a *binary predicate*.
- This terminology is used in much of the available C++ documentation and literature.



Lambda Expressions

- ➊ A lambda expression is a compact way of creating a function object without having to write a class declaration.
- ➋ It is an expression that contains only the logic of the object's operator() member function.
- ➌ When the compiler encounters a lambda expression, it automatically generates a function object in memory, using the code that you provide in the lambda expression for the operator() member function.



Lambda Expressions

- General format:

`[](parameter list) { function body }`

- The [] is known as the lambda introducer. It marks the beginning of a lambda expression.
- parameter list* is a list of parameter declarations for the function object's operator() member function.
- function body* is the code that should be the body of the object's operator() member function.

Lambda Expressions

- Example: a lambda expression for a function object that computes the sum of two integers:

```
[](int a, int b) { return x + y; }
```

Lambda Expressions

- Example: a lambda expression for a function object that determines whether an integer is even is:

```
[](int x) { return x % 2 == 0; }
```

Lambda Expressions

- Example: a lambda expression for a function object that takes an integer as input and prints the square of that integer:

```
[](int a) { cout << a * a << " ";
```

Lambda Expressions

- When you call a lambda expression, you write a list of arguments, enclosed in parentheses, right after the expression.
- For example, the following code snippet displays 7, which is the sum of the variables x and y:

```
int x = 2;  
int y = 5;  
cout << [](int a, int b) {return a + b;}(x, y) << endl;
```

Lambda Expressions

- The following code segment counts the even numbers in a vector:

```
// Create a vector of ints.  
vector<int> v = { 1, 2, 3, 4, 5, 6, 7, 8 };  
  
// Get the number of elements that are even.  
int evenNums = count_if(v.begin(), v.end(), [](int x) {return x % 2 == 0;});  
  
// Display the results.  
cout << "The vector contains " << evenNums << " even numbers.\n";
```

Lambda Expressions

- Because lambda expressions generate function objects, you can assign a lambda expression to a variable and then call it through the variable's name:

```
auto sum = [](int a, int b) {return a + b;};  
int x = 2;  
int y = 5;  
int z = sum(x, y);
```

Program 17-37

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4 using namespace std;
5
6 int main()
7 {
8     // Create a vector of ints.
9     vector<int> v = { 1, 2, 3, 4, 5, 6, 7, 8 };
10
11    // Use a lambda expression to create a function object.
12    auto isEven = [] (int x) { return x % 2 == 0; };
13
14    // Get the number of elements that even.
15    int evenNums = count_if(v.begin(), v.end(), isEven);
16
17    // Display the results.
18    cout << "The vector contains " << evenNums << " even numbers.\n";
19    return 0;
20 }
```

Program Output

The vector contains 4 even numbers.



Functional Classes in the STL

- The STL library defines a number of classes that you can instantiate to create function objects in your program.
- To use these classes, you must `#include` the `<functional>` header file.
- Table 17-15 in your textbook lists a few of the functional classes:

Table 17-15 STL Function Object Classes

Functional Class	Description
<code>less<T></code>	<code>less<T>()</code> (<code>T a, T b</code>) is true if and only if $a < b$
<code>less_equal<T></code>	<code>less_equal()</code> (<code>T a, T b</code>) is true if and only if $a \leq b$
<code>greater<T></code>	<code>greater<T>()</code> (<code>T a, T b</code>) is true if and only if $a > b$
<code>greater_equal<T></code>	<code>greater_equal<T>()</code> (<code>T a, T b</code>) is true if and only if $a \geq b$