




## K Nearest Neighbor Classification

plus similarity and distance calculations

# K Nearest Neighbor Classification

- K nearest neighbors is a simple algorithm that stores all available cases and classifies new cases based on a similarity measure (e.g., distance functions).
- A case is classified by a majority vote of its neighbors, with the case being assigned to the class most common amongst its K nearest neighbors measured by a distance function.



# Type of Data in Clustering Analysis

- Interval-scaled variables
- Binary variables
- Nominal, and ordinal variables
- Variables of mixed types
- Text
- Temporal

# Standardize Numeric Data

- Standardize data

- Calculate the mean absolute deviation:

$$s_f = \frac{1}{n} (|x_{1f} - m_f| + |x_{2f} - m_f| + \dots + |x_{nf} - m_f|)$$

Where

$$m_f = \frac{1}{n} (x_{1f} + x_{2f} + \dots + x_{nf}).$$

- Calculate the standardized measurement (z-score)

$$z_{if} = \frac{x_{if} - m_f}{s_f}$$

- Normalizing data

$$z_{if} = \frac{x_{if} - m_f}{\sigma_f}$$

# Similarity/Dissimilarity Between Objects

- Distances are normally used to measure the similarity or dissimilarity between two data objects
- Some popular ones include: *Minkowski distance*:

$$d(i, j) = \sqrt[q]{(|x_{i1} - x_{j1}|^q + |x_{i2} - x_{j2}|^q + \dots + |x_{ip} - x_{jp}|^q)}$$

Where  $i = (x_{i1}, x_{i2}, \dots, x_{ip})$  and  $j = (x_{j1}, x_{j2}, \dots, x_{jp})$  are two  $p$ -dimensional data objects, and  $q$  is a positive integer

- If  $q = 1$ ,  $d$  is *Manhattan distance*

$$d(i, j) = |x_{i1} - x_{j1}| + |x_{i2} - x_{j2}| + \dots + |x_{ip} - x_{jp}|$$

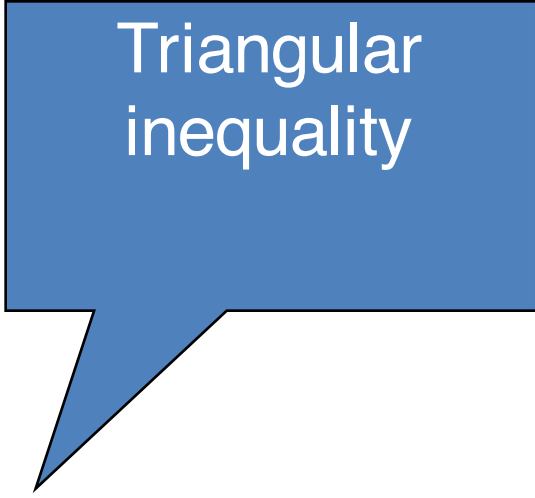
# Similarity/Dissimilarity Between Objects

If  $q = 2$ ,  $d$  is Euclidean distance:

$$d(i, j) = \sqrt{(|x_{i_1} - x_{j_1}|^2 + |x_{i_2} - x_{j_2}|^2 + \dots + |x_{i_p} - x_{j_p}|^2)}$$

- Properties

- $d(i, j) \geq 0$
- $d(i, i) = 0$
- $d(i, j) = d(j, i)$
- $d(i, j) \leq d(i, k) + d(k, j)$



Triangular inequality

# Other Similarity/Distance measures

- **Sets as vectors:** measure similarity by the cosine distance.

$$x_i = [x_{i1}, x_{i2}, \dots, x_{ip}]$$

$$x_j = [x_{j1}, x_{j2}, \dots, x_{jp}]$$

$$\cos(x_i, x_j) = \frac{x_i \bullet x_j}{|x_i| * |x_j|} = \hat{x}_i \bullet \hat{x}_j$$

# Similarity/Dissimilarity for Binary Data

- **Symmetric attribute:** both states are equally valuable, carrying the same weight, e.g., gender 0, 1  $\rightarrow$  coding each state to be 0 or 1 arbitrarily
- **Asymmetric attribute:** outcomes of the states are not equally important, e.g., outcome of a medical test  $\rightarrow$  positive or negative
  - assign the more important outcome to value 1
  - assign the less important outcome to value 0
  - e.g., 1 : HIV positive, 0: HIV negative. The agreement of two patients having “1”s for this attribute is more significant than an agreement of “0”s.



# Similarity/Dissimilarity for Binary Data

- A contingency table for binary data

		Object $j$		
		1	0	$sum$
Object $i$	1	$a$	$b$	$a+b$
	0	$c$	$d$	$c+d$
	$sum$	$a+c$	$b+d$	$p$

- Simple matching coefficient (if the binary variable is symmetric):

$$d(i, j) = \frac{b+c}{a+b+c+d}$$

- Jaccard coefficient (if the binary variable is asymmetric):

$$d(i, j) = \frac{b+c}{a+b+c}$$

# Dissimilarity between Binary Variables

- Example

Name	Gender	Fever	Cough	Test-1	Test-2	Test-3	Test-4
Jack	M	Y	N	P	N	N	N
Mary	F	Y	N	P	N	P	N
Jim	M	Y	P	N	N	N	N

- gender is a symmetric attribute
- the remaining attributes are asymmetric binary
- let the values Y and P be set to 1, and the value N be set to 0

		Jack	
Mary		1	0
	1		
	0		

		Jack	
Jim		1	0
	1		
	0		

		Jim	
Mary		1	0
	1		
	0		

# Dissimilarity between Binary Variables

Jaccard  
coefficient

$$d(\text{jack}, \text{mary}) = \frac{0 + 1}{2 + 0 + 1} = 0.33$$

$$d(\text{jack}, \text{jim}) = \frac{1 + 1}{1 + 1 + 1} = 0.67$$

$$d(\text{jim}, \text{mary}) = \frac{1 + 2}{1 + 1 + 2} = 0.75$$

Gender is not yet included in the computation

		Jack	
Mary		1	0
	1		
	0		

		Jack	
Jim		1	0
	1		
	0		

		Jim	
Mary		1	0
	1		
	0		

# Nominal Attributes

- A generalization of the binary attribute in that it can take more than 2 states, e.g., red, yellow, blue, green
- Method 1: Simple matching
  - $m$ : # of matches,  $p$ : total # of variables
$$d(i, j) = \frac{p - m}{p}$$
- Method 2: use a large number of binary attributes
  - creating a new binary variable for each of the  $M$  nominal states

# Ordinal Attributes

- An ordinal attribute can be discrete or continuous
- order is important, e.g., rank
- Can be treated like interval-scaled
  - replacing  $x_{if}$  by their rank  $r_{if} \in \{1, \dots, M_f\}$
  - map the range of each attribute onto  $[0, 1]$  by replacing  $i$ -th object in the  $f$ -th attribute by
$$z_{if} = \frac{r_{if} - 1}{M_f - 1}$$
  - compute the dissimilarity using methods for interval-scaled attributes

# Attributes of Mixed Types

- A database may contain different types of attributes
  - symmetric binary, asymmetric binary, nominal, ordinal, and interval.
- How to combine the dissimilarity/distance from data of a mixture of types?

# Attributes of Mixed Types

- Use weighted formula to combine their effects.

- Feature value missing, or asymmetric binary with

$$x_{if} = x_{jf} = 0 \quad \rightarrow \quad \delta_{ij}(f) = 0$$

- Otherwise  $\rightarrow \delta_{ij}(f) = 1$

- feature is interval-based: use the normalized distance ( $\delta_{ij}^{(f)}$  is weight on feature  $f$ )

$$d(i, j) = \frac{\sum_{f=1}^p \delta_{ij}^{(f)} d_{ij}^{(f)}}{\sum_{f=1}^p \delta_{ij}^{(f)}}$$

- feature is ordinal

- compute ranks  $r_{if}$  and
- and treat  $z_{if}$  as interval-scaled

$$z_{if} = \frac{r_{if} - 1}{M_f - 1}$$

# Practice Question

- Compute the distance between (obj1, obj2),

	Gender	Age	Heart Rate	Fever	Cough	Category
Obj1	F	18	120	N	N	Severe-1
Obj2	M	36	89	N	N	Normal

For Age:  $m=42$ ,  $s=3.5$ ,

For heart rate:  $m=95$ ,  $s=10$

Possible values for Category include : Normal, Severe-1, Severe-2, Dying

For simplicity in demonstration, use Manhattan distance for interval data.



# Other Similarity/Distance measures

- Measure distance between words/address/query, or between DNA sequences by edit distance
  - Given two strings  $S_1$  and  $S_2$ , the minimum number of operations to convert one to the other
  - Operations are typically character-level
    - Insert, Delete, Replace, (Transposition)
  - E.g., the edit distance from **dof** to **dog** is 1
    - From **cat** to **act** is 2 (Just 1 with transpose.)
    - from **cat** to **dog** is 3.
- Generally computed by dynamic programming.

18

# Edit distance

- Given two strings  $S_1$  and  $S_2$ , the **minimum** number of operations to convert one to the other
- Operations are typically character-level
  - Insert, Delete, Replace, (Transposition)
- E.g., the edit distance from **dof** to **dog** is 1
  - From **cat** to **act** is 2 (Just 1 with transpose.)
  - from **cat** to **dog** is 3.
- Generally found by dynamic programming.

# Edit Distance

EDITDISTANCE( $s_1, s_2$ )

```
1  int  $m[|s_1|, |s_2|] = 0$ 
2  for  $i \leftarrow 1$  to  $|s_1|$ 
3  do  $m[i, 0] = i$ 
4  for  $j \leftarrow 1$  to  $|s_2|$ 
5  do  $m[0, j] = j$ 
6  for  $i \leftarrow 1$  to  $|s_1|$ 
7  do for  $j \leftarrow 1$  to  $|s_2|$ 
8      do  $m[i, j] = \min\{m[i - 1, j - 1] + \text{if } (s_1[i] = s_2[j]) \text{ then } 0 \text{ else } 1, \text{ fi},$ 
9           $m[i - 1, j] + 1,$ 
10          $m[i, j - 1] + 1\}$ 
11 return  $m[|s_1|, |s_2|]$ 
```

Figure 3.5 Dynamic programming algorithm for computing the edit distance between strings  $s_1$  and  $s_2$ .

# An Example

		f	a	s	t
	<hr/> 0	<hr/> 1 1	<hr/> 2 2	<hr/> 3 3	<hr/> 4 4
c	<hr/> 1 1	<hr/> 1 2 2 1	<hr/> 2 3 2 2	<hr/> 3 4 3 3	<hr/> 4 5 4 4
a	<hr/> 2 2	<hr/> 2 2 3 2	<hr/> 1 3 3 1	<hr/> 3 4 2 2	<hr/> 4 5 3 3
t	<hr/> 3 3	<hr/> 3 3 4 3	<hr/> 3 2 4 2	<hr/> 2 3 3 2	<hr/> 2 4 3 2
s	<hr/> 4 4	<hr/> 4 4 5 4	<hr/> 4 3 5 3	<hr/> 2 3 4 2	<hr/> 3 3 3 3

# Weighted edit distance

- As above, but the weight of an operation depends on the character(s) involved
  - Meant to capture OCR or keyboard errors, e.g. *m* more likely to be mis-typed as *n* than as *q*
  - Therefore, replacing *m* by *n* is a smaller edit distance than by *q*
  - This may be formulated as a probability model
- Requires weight matrix as input
- Modify dynamic programming to handle weights

# Finding nearest neighbors efficiently


- Simplest way of finding nearest neighbor: linear scan of the data
  - ◆ Classification takes time proportional to the product of the number of instances in training and test sets
- Nearest-neighbor search can be done more efficiently using appropriate data structures
- Two methods that represent training data in a tree structure:

*kD-trees* and *ball trees*



*(k-dimensional tree)*

# $k$ D-tree example

- Binary tree in which every node is a  $k$ -dimensional point.
- Every non-leaf node can be thought of as implicitly generating a splitting hyperplane that divides the space into two parts, known as half-spaces.



# kD-tree example



The hyperplane direction is chosen in the following way:  
every node in the tree is associated with one of the  $k$ -dimensions,  
with the hyperplane perpendicular to that dimension's axis.





# kD-tree example


```
function kdtree (list of points pointList, int depth)
{
    // Select axis based on depth so that axis cycles through all valid values
    var int axis := depth mod k;

    // Sort point list and choose median as pivot element
    select median by axis from pointList;



    // Create node and construct subtree
    node.location := median;
    node.leftChild := kdtree(points in pointList before median, depth+1);
    node.rightChild := kdtree(points in pointList after median, depth+1);
    return node;
}
```



# Using kD-trees: example (1)





# Using kD-trees: example (2)




Start at A, then proceed in depth-first search (maintain a stack of parent-nodes if using a singly-linked tree). Set best estimate to A's distance. Then examine left child node

# Using kD-trees: example (3)



Calculate B's distance and compare against best estimate  
- It is smaller distance, so update best estimate. Examine children (left then right)


# Using kD-trees: example (4)



Went left of split

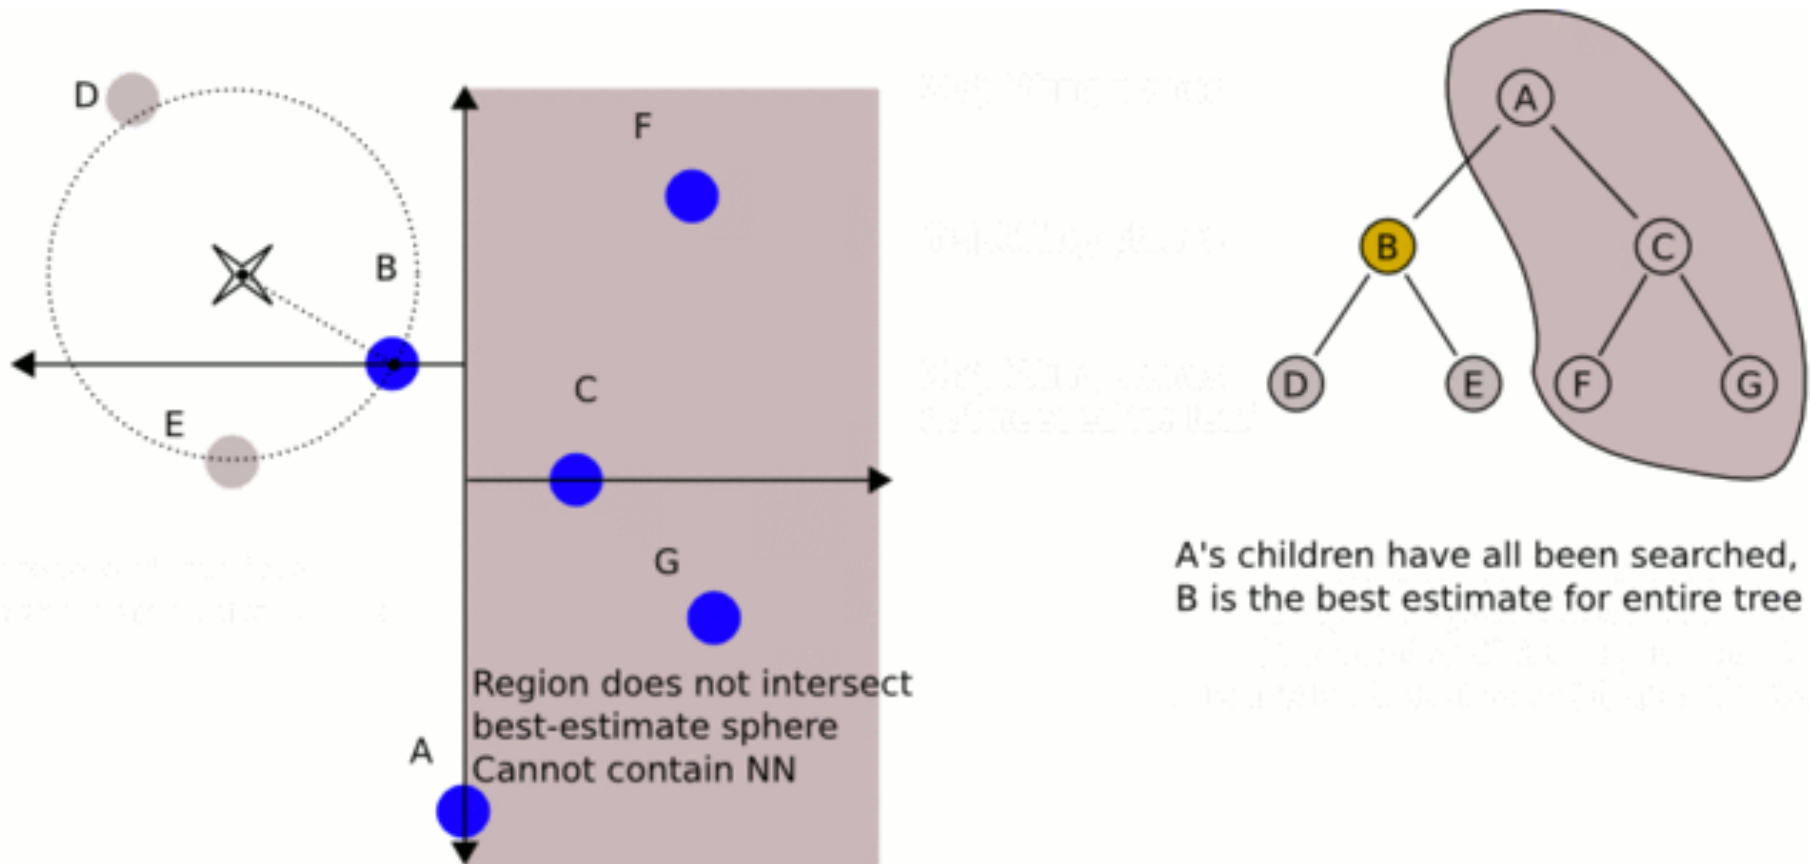
Pruning Allowed

We know that  
C is closer than B



D & E Discarded as B  
(already visited) is closer.  
B is the best estimate for B's sub-branch  
Proceed back to parent node

# Using kD-trees: example (5)



# More on *kD*-trees

- Complexity depends on depth of tree, given by logarithm of number of nodes
- Amount of backtracking required depends on quality of tree (“square” vs. “skinny” nodes)
- How to build a good tree? Need to find good split point and split direction
  - ◆ Split direction: direction with greatest variance
  - ◆ Split point: median value along that direction
- Using value closest to mean (rather than median) can be better if data is skewed
- Can apply this recursively

# Building trees incrementally



- Big advantage of instance-based learning: classifier can be updated incrementally
  - ◆ Just add new training instance!
- Can we do the same with  $k$ D-trees?
- Heuristic strategy:
  - ◆ Find leaf node containing new instance
  - ◆ Place instance into leaf if leaf is empty
  - ◆ Otherwise, split leaf according to the longest dimension (to preserve squareness)
- Tree should be re-built occasionally (i.e. if depth grows to twice the optimum depth)



# Ball trees


- Problem in  $k$ D-trees: corners
- Observation: no need to make sure that regions don't overlap
- Can use balls (hyperspheres) instead of hyperrectangles
  - ◆ A *ball tree* organizes the data into a tree of  $k$ -dimensional hyperspheres
  - ◆ Normally allows for a better fit to the data and thus more efficient search

# Ball tree example



# Using ball trees

- Nearest-neighbor search is done using the same backtracking strategy as in *kD*-trees
- Ball can be ruled out from consideration if: distance from target to ball's center exceeds ball's radius plus current upper bound



# Building ball trees

- Ball trees are built top down (like *kD*-trees)
- Basic problem: splitting a ball into two

```
function construct_balltree is
  input:
    D, an array of data points
  output:
    B, the root of a constructed ball tree
  if a single point remains then
    create a leaf B containing the single point in D
    return B
  else
    let c be the dimension of greatest spread
    let L,R be the sets of points lying to the left and right of the median along dimension c
    create B with two children:
      B.pivot = c
      B.child1 = construct_balltree(L),
      B.child2 = construct_balltree(R)
    return B
  end if
end function
```

# Nearest Neighbor with Ball Tree

- At each node  $B$ , it may perform one of three operations, before finally returning an updated version of the priority queue:
  - If the distance from the test point  $t$  to the current node  $B$  is greater than the furthest point in  $Q$  (*a maximum first heap*), ignore  $B$  and return  $Q$ .
  - If  $B$  is a leaf node, scan through every point enumerated in  $B$  and update the nearest-neighbor queue appropriately. Return the updated queue.
  - If  $B$  is an internal node, call the algorithm recursively on  $B$ 's two children, searching the child whose center is closer to  $t$  first. Return the queue after each of these calls has updated it in turn.

**function knn\_search is**

**input:**

t, the target point for the query

k, the number of nearest neighbors of t to search for

Q, max-first priority queue containing at most k points

B, a node, or ball, in the tree

**output:**

Q, containing the k nearest neighbors from within B

**if** distance(t, B.pivot)  $\geq$  distance(t, Q.first) **then**

**return** Q unchanged

**else if** B is a leaf node **then**

**for each** point p in B **do**

**if** distance(t, p) < distance(t, Q.first) **then**

            add p to Q

**if** size(Q) > k **then**

                remove the furthest neighbor from Q

**end if**

**end if**

**repeat**

**else**

    let child1 be the child node closest to t

    let child2 be the child node furthest from t

    knn\_search(t, k, Q, child1)

    knn\_search(t, k, Q, child2)

**end if**

**end function**<sup>[2]</sup>

# Discussion of nearest-neighbor learning

- Often very accurate
- Assumes all attributes are equally important
  - Remedy: attribute selection or weights
- Possible remedies against noisy instances:
  - Take a majority vote over the  $k$  nearest neighbors
  - Removing noisy instances from dataset (difficult!)
- Statisticians have used  $k$ -NN since early 1950s
  - If  $n \rightarrow \infty$  and  $k/n \rightarrow 0$ , error approaches minimum
- $k$ D-trees become inefficient when number of attributes is too large (approximately  $> 10$ )
- Ball trees (which are instances of *metric trees*) work well in higher-dimensional spaces