

# Data Mining



## Bayesian Classifier

# A Quick Review of Probability

- The Axioms of Probability
  - $0 \leq P(A) \leq 1$
  - $P(\text{True}) = 1$
  - $P(\text{False}) = 0$
  - $P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$

# Theorems from the Axioms

- $0 \leq P(A) \leq 1$ ,  $P(\text{True}) = 1$ ,  $P(\text{False}) = 0$
- $P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$

From these we can prove:

$$P(\text{not } A) = P(\sim A) = 1 - P(A)$$

$$P(A) = P(A \text{ and } B) + P(A \text{ and } \sim B)$$

# Multivalued Random Variables

- Suppose  $A$  can take on more than 2 values
- $A$  is a *random variable with arity  $k$*  if it can take on exactly one value out of  $\{v_1, v_2, \dots, v_k\}$
- Thus...

$$P(A = v_i \text{ and } A = v_j) = 0 \text{ if } i \neq j$$

$$P(A = v_1 \text{ or } A = v_2 \cdots \text{ or } A = v_k) = 1$$

# An easy fact about Multivalued Random Variables

- Using the axioms of probability...

$$0 \leq P(A) \leq 1, P(\text{True}) = 1, P(\text{False}) = 0$$

$$P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$$

- And assuming that A obeys...

$$P(A = v_i \wedge A = v_j) = 0 \text{ if } i \neq j$$

$$P(A = v_1 \vee A = v_2 \vee \dots \vee A = v_k) = 1$$

- It's easy to prove that

$$P(A = v_1 \vee A = v_2 \vee \dots \vee A = v_i) = \sum_{j=1}^i P(A = v_j)$$

- And thus we can prove

$$\sum_{j=1}^k P(A = v_j) = 1$$

# Another fact about Multivalued Random Variables:

- Using the axioms of probability...

$$0 \leq P(A) \leq 1, P(\text{True}) = 1, P(\text{False}) = 0$$

$$P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$$

- And assuming that A obeys...

$$P(A = v_i \text{ and } A = v_j) = 0 \text{ if } i \neq j$$

$$P(A = v_1 \text{ or } A = v_2 \text{ or } A = v_k) = 1$$

- It can be proved that

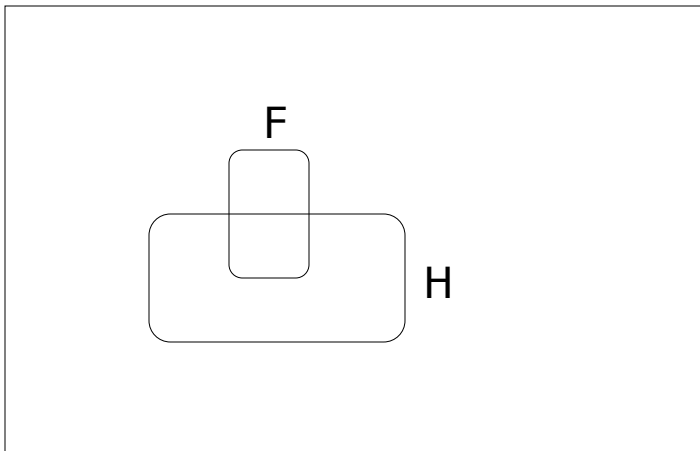
$$P(B \text{ and } [A = v_1 \text{ or } A = v_2 \text{ or } A = v_i]) = \sum_{j=1}^i P(B \text{ and } (A = v_j))$$

- And thus we can prove

$$P(B) = \sum_{j=1}^k P(B \text{ and } A = v_j)$$

# Conditional Probability

- $P(A | B)$  = Fraction of worlds in which B is true that also have A true



H = “Have a headache”  
F = “Coming down with Flu”

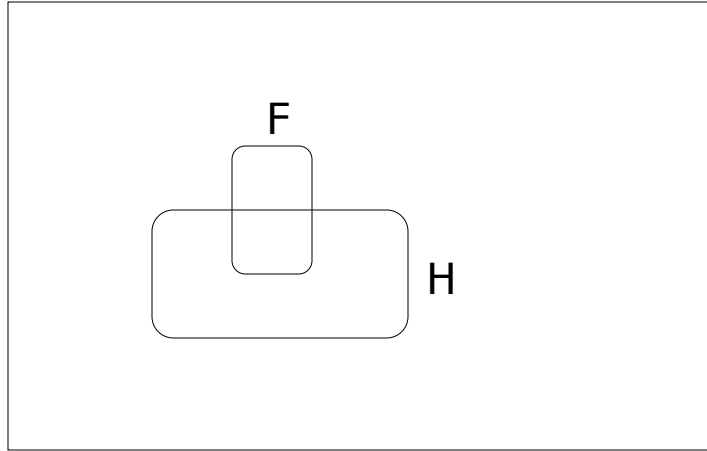
$$P(H) = 1/10$$

$$P(F) = 1/40$$

$$P(H|F) = 1/2$$

“Headaches are rare and flu is rarer, but if you’re coming down with flu there’s a 50-50 chance you’ll have a headache.”

# Conditional Probability



H = “Have a headache”  
F = “Coming down with Flu”

$$\begin{aligned}P(H) &= 1/10 \\P(F) &= 1/40 \\P(H|F) &= 1/2\end{aligned}$$

$P(H|F)$  = Fraction of flu-inflicted worlds in which you have a headache

$$= \frac{\text{\#worlds with flu and headache}}{\text{\#worlds with flu}}$$

$$= \frac{\text{Area of “H and F” region}}{\text{Area of “F” region}}$$

$$= \frac{P(H \text{ and } F)}{P(F)}$$



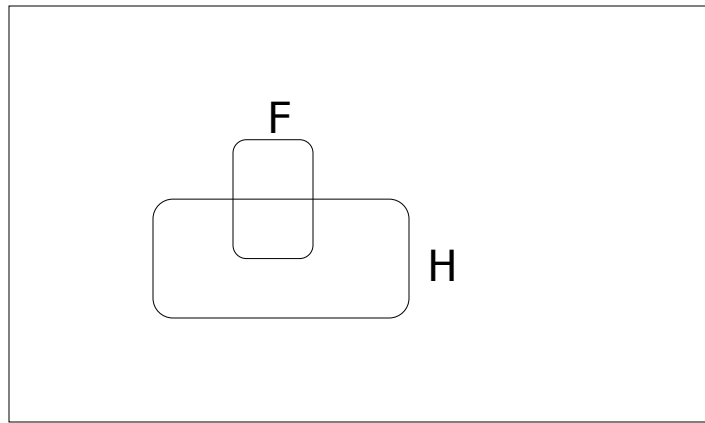
# Definition of Conditional Probability

$$P(A|B) = \frac{P(A \text{ and } B)}{P(B)}$$

## Corollary: The Chain Rule

$$P(A \text{ and } B) = P(A|B) P(B)$$

# Probabilistic Inference



H = “Have a headache”

F = “Coming down with Flu”

$$P(H) = 1/10$$

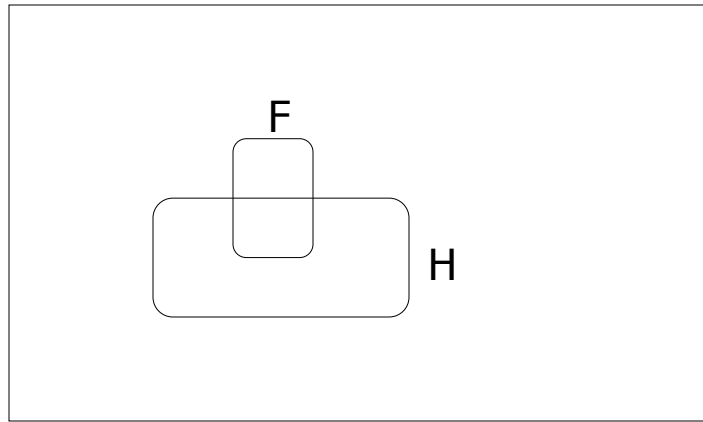
$$P(F) = 1/40$$

$$P(H|F) = 1/2$$

One day you wake up with a headache. You think: “Drat! 50% of flus are associated with headaches so I must have a 50-50 chance of coming down with flu”

Is this reasoning correct?

# Probabilistic Inference



H = “Have a headache”  
F = “Coming down with Flu”

$$P(H) = 1/10$$

$$P(F) = 1/40$$

$$P(H|F) = 1/2$$

$$P(F \text{ and } H) = \dots$$

$$P(F|H) = \dots$$

# What we just did...is the Bayes Rule

$$P(A|B) = \frac{P(A \text{ and } B)}{P(B)} = \frac{P(B|A) P(A)}{P(B)}$$

**Bayes, Thomas (1763)** An essay towards solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society of London*, **53:370-418**



# More General Forms of Bayes Rule

$$P(A|B) = \frac{P(B|A)P(A)}{P(B|A)P(A) + P(B|\sim A)P(\sim A)}$$

$$P(A = v_i | B) = \frac{P(B | A = v_i)P(A = v_i)}{\sum_{k=1}^{n_A} P(B | A = v_k)P(A = v_k)}$$

# Useful Easy-to-prove facts

$$P(A | B) + P(\neg A | B) = 1$$

$$\sum_{k=1}^{n_A} P(A = v_k | B) = 1$$

# The Joint Distribution

Recipe for making a joint distribution  
of  $M$  variables:

*Example: Boolean  
variables  $A, B, C$*

# The Joint Distribution

*Example: Boolean variables A, B, C*

Recipe for making a joint distribution of M variables:

1. Make a truth table listing all combinations of values of your variables (if there are M Boolean variables then the table will have  $2^M$  rows).

A	B	C
0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0
1	1	1



# The Joint Distribution

Recipe for making a joint distribution of  $M$  variables:

1. Make a truth table listing all combinations of values of your variables (if there are  $M$  Boolean variables then the table will have  $2^M$  rows).
2. For each combination of values, say how probable it is.

*Example: Boolean variables  $A, B, C$*

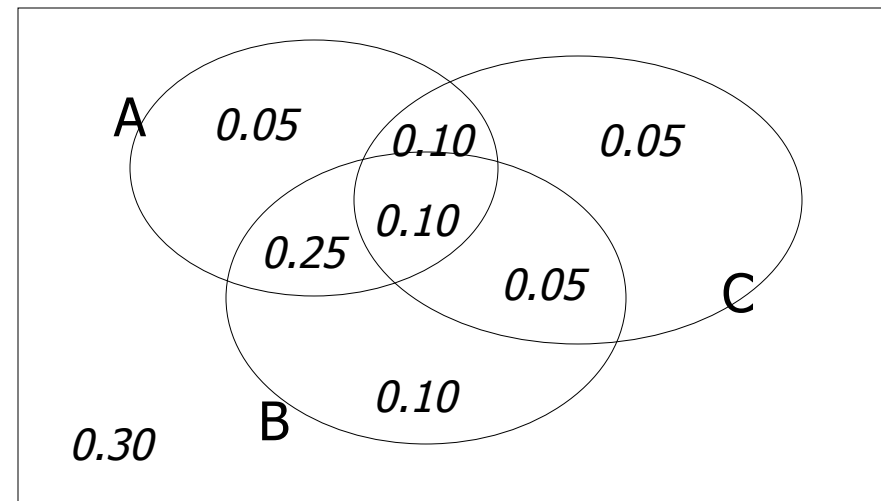
<b>A</b>	<b>B</b>	<b>C</b>	<b>Prob</b>
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

# The Joint Distribution

Recipe for making a joint distribution of  $M$  variables:

1. Make a truth table listing all combinations of values of your variables (if there are  $M$  Boolean variables then the table will have  $2^M$  rows).
2. For each combination of values, say how probable it is.
3. If you subscribe to the axioms of probability, those numbers must sum to 1.

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10



# Using the Joint

gender	hours_worked	wealth		
Female	v0:40.5-	poor	0.253122	
		rich	0.0245895	
	v1:40.5+	poor	0.0421768	
		rich	0.0116293	
Male	v0:40.5-	poor	0.331313	
		rich	0.0971295	
	v1:40.5+	poor	0.134106	
		rich	0.105933	

Once you have the JD you can ask for the probability of any logical expression involving your attribute

$$P(E) = \sum_{\text{rows matching } E} P(\text{row})$$









# Using the Joint

gender	hours_worked	wealth		
Female	v0:40.5-	poor	0.253122	<div></div>
		rich	0.0245895	<div></div>
	v1:40.5+	poor	0.0421768	<div></div>
		rich	0.0116293	<div></div>
Male	v0:40.5-	poor	0.331313	<div></div>
		rich	0.0971295	<div></div>
	v1:40.5+	poor	0.134106	<div></div>
		rich	0.105933	<div></div>

$$P(\text{Poor Male}) = 0.4654$$

$$P(E) = \sum_{\text{rows matching } E} P(\text{row})$$








# Using the Joint

gender	hours_worked	wealth	
Female	v0:40.5-	poor	0.253122 
		rich	0.0245895 
	v1:40.5+	poor	0.0421768 
		rich	0.0116293 
Male	v0:40.5-	poor	0.331313 
		rich	0.0971295 
	v1:40.5+	poor	0.134106 
		rich	0.105933 

$$P(\text{Poor}) = 0.7604$$

$$P(E) = \sum_{\text{rows matching } E} P(\text{row})$$

# Inference with the Joint

gender	hours_worked	wealth		
Female	v0:40.5-	poor	0.253122	
		rich	0.0245895	
	v1:40.5+	poor	0.0421768	
		rich	0.0116293	
Male	v0:40.5-	poor	0.331313	
		rich	0.0971295	
	v1:40.5+	poor	0.134106	
		rich	0.105933	

$$P(E_1 | E_2) = \frac{P(E_1 \text{ and } E_2)}{P(E_2)} = \frac{\sum_{\text{rows matching } E_1 \text{ and } E_2} P(\text{row})}{\sum_{\text{rows matching } E_2} P(\text{row})}$$

# Inference with the Joint

gender	hours_worked	wealth	
Female	v0:40.5-	poor	0.253122
		rich	0.0245895
	v1:40.5+	poor	0.0421768
		rich	0.0116293
Male	v0:40.5-	poor	0.331313
		rich	0.0971295
	v1:40.5+	poor	0.134106
		rich	0.105933

$$P(E_1 | E_2) = \frac{P(E_1 \text{ and } E_2)}{P(E_2)} = \frac{\sum_{\text{rows matching } E_1 \text{ and } E_2} P(\text{row})}{\sum_{\text{rows matching } E_2} P(\text{row})}$$

$$P(\text{Male} | \text{Poor}) = 0.4654 / 0.7604 = 0.612$$

# Inference is a big deal

- I've got this evidence. What's the chance that this conclusion is true?
  - I've got a sore neck: how likely am I to have meningitis?
  - I see my lights are out and it's 9pm. What's the chance my spouse is already asleep?
- There's a thriving set of industries growing based around Bayesian Inference. Highlights are: Medicine, Pharma, Help Desk Support, Engine Fault Diagnosis



# Where do Joint Distributions come from?

- Idea One: Expert Humans
- Idea Two: Simpler probabilistic facts and some algebra

Example: Suppose you knew

$$\begin{array}{ll} P(A) = 0.7 & P(C|A \text{ and } B) = 0.1 \\ & P(C|A \text{ and } \sim B) = 0.2 \\ P(B|A) = 0.2 & P(C|\sim A \text{ and } B) = 0.3 \\ P(B|\sim A) = 0.1 & P(C|\sim A \text{ and } \sim B) = 0.1 \end{array}$$

Then you can automatically compute the JD using the chain rule

$$P(A=x \text{ and } B=y \text{ and } C=z) = P(C=z|A=x \text{ and } B=y) P(B=y|A=x) P(A=x)$$

What is  $P(A, B, \sim C)$ ?

# Where do Joint Distributions come from?

- Idea Three: Learn them from data!

Prepare to see one of the most impressive learning algorithms you'll come across in the entire course....

# Learning a joint distribution

Build a JD table for your attributes in which the probabilities are unspecified

A	B	C	Prob
0	0	0	?
0	0	1	?
0	1	0	?
0	1	1	?
1	0	0	?
1	0	1	?
1	1	0	?
1	1	1	?

Fraction of all records in which  
A and B are True but C is False






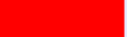


The fill in each row with

$$\hat{P}(\text{row}) = \frac{\text{records matching row}}{\text{total number of records}}$$

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

# Example of Learning a Joint

- This Joint was obtained by learning from three attributes in the UCI “Adult” Census Database [Kohavi 1995]

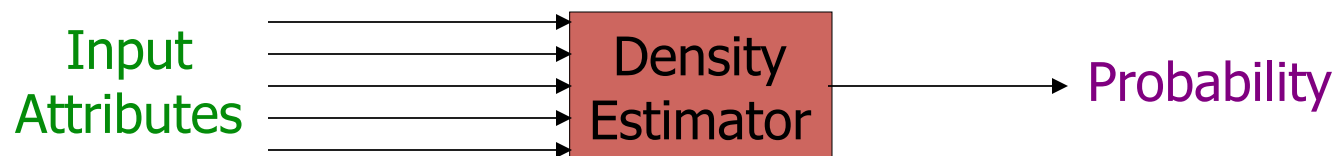
gender	hours_worked	wealth		
Female	v0:40.5-	poor	0.253122	
		rich	0.0245895	
	v1:40.5+	poor	0.0421768	
		rich	0.0116293	
Male	v0:40.5-	poor	0.331313	
		rich	0.0971295	
	v1:40.5+	poor	0.134106	
		rich	0.105933	

# Where are we?

- We have recalled the fundamentals of probability
- We have become content with what JDs are and how to use them
- We know how to learn JDs from data.

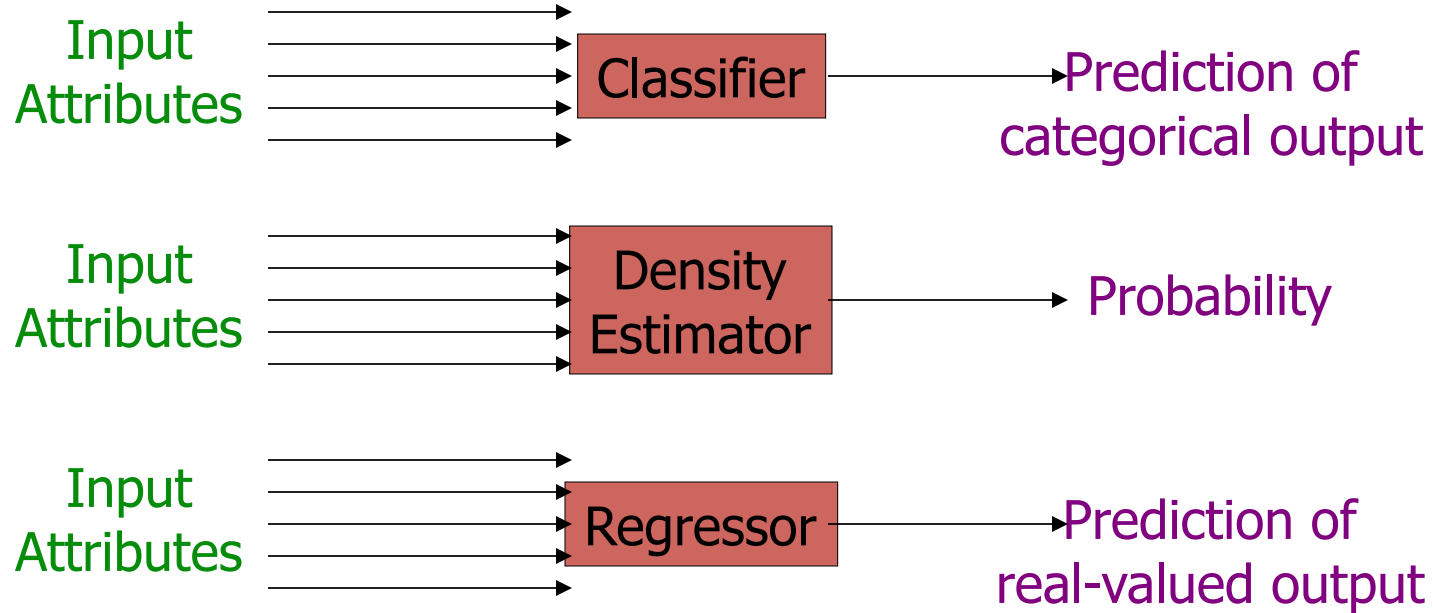
# Density Estimation

- Our Joint Distribution learner is our first example of something called Density Estimation
- A Density Estimator learns a mapping from a set of attributes to a Probability



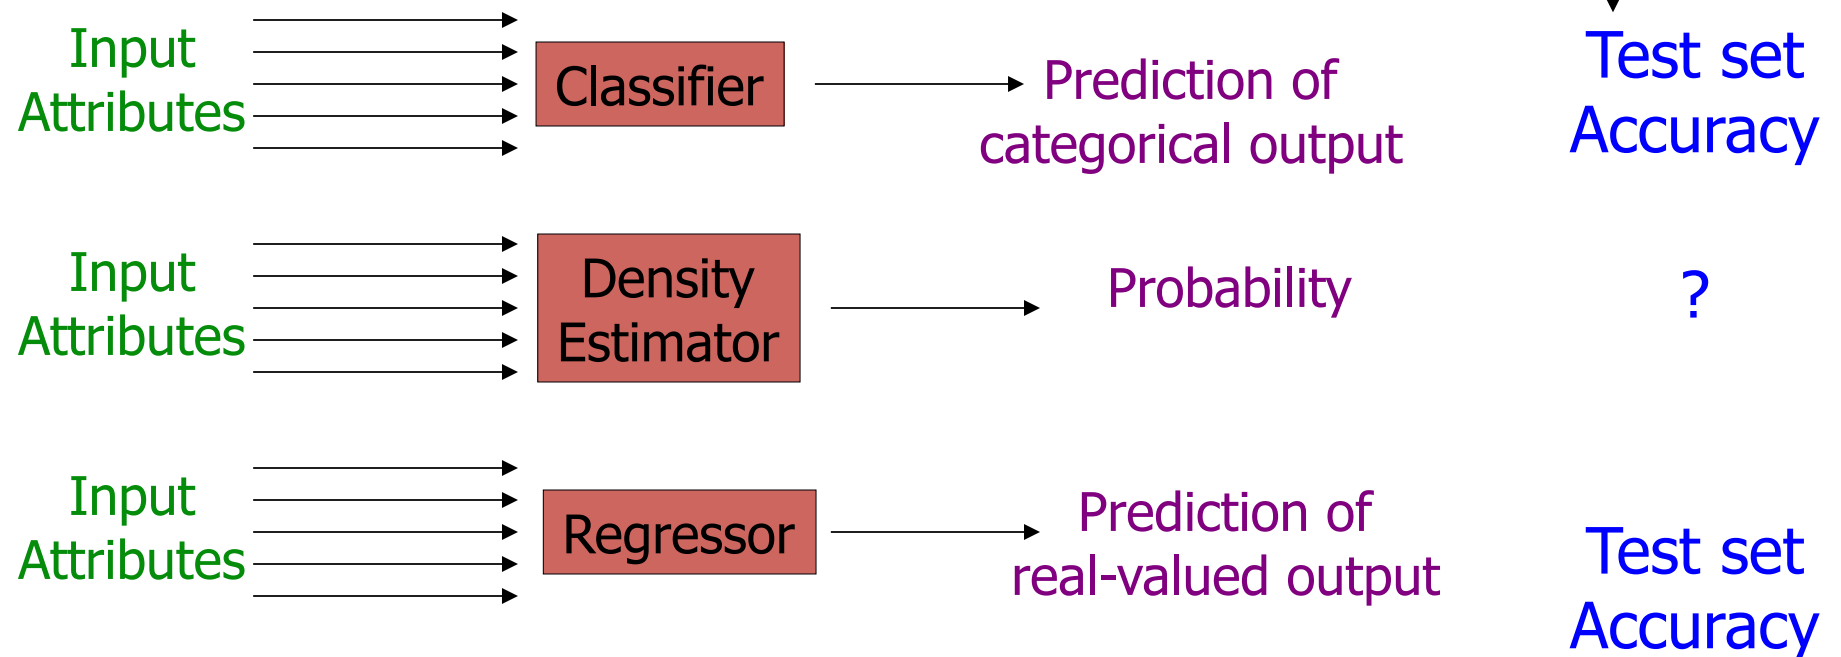
# Density Estimation

- Compare it against the two other major kinds of models:



# Evaluating Density Estimation

Test-set criterion for estimating performance on future data





# Evaluating a density estimator

- Given a record  $\mathbf{x}$ , a density estimator  $M$  can tell you how likely the record is:

$$\hat{P}(\mathbf{x}|M)$$

- Given a dataset with  $R$  records, a density estimator can tell you how likely the dataset is:

(Under the assumption that all records were independently generated from the Density Estimator's JD)

$$\hat{P}(\text{dataset} | M) = \hat{P}(\mathbf{x}_1 \text{ and } \mathbf{x}_2 \dots \text{ and } \mathbf{x}_R | M) = \prod_{k=1}^R \hat{P}(\mathbf{x}_k | M)$$

# Revisit the Miles Per Gallon dataset

192  
Training  
Set  
Records

mpg	modelyear	maker
good	75to78	asia
bad	70to74	america
bad	75to78	europa
bad	70to74	america
bad	70to74	america
bad	70to74	asia
bad	70to74	asia
bad	75to78	america
:	:	:
:	:	:
:	:	:
bad	70to74	america
good	79to83	america
bad	75to78	america
good	79to83	america
bad	75to78	america
good	79to83	america
good	79to83	america
bad	70to74	america
good	75to78	europa
bad	75to78	europa

# the Miles Per Gallon dataset

192  
Training  
Set  
Records

mpg	modelyear	maker
good	75to78	asia
bad	70to74	america
bad	75to78	europe
bad	70to74	america
bad	70to74	america
bad	70to74	asia
bad	70to74	asia
bad	75to78	america
:	:	:
:	:	:
:	:	:
bad	70to74	america
good	79to83	america
bad	75to78	america
good	79to83	america
bad	75to78	america
good	79to83	america
good	79to83	america
bad	70to74	america
good	75to78	europe
bad	75to78	europe

mpg	modelyear	maker		
bad	70to74	america	0.27551	<div></div>
		asia	0.0255102	<div></div>
		europe	0.0153061	<div></div>
	75to77	america	0.153061	<div></div>
		asia	0.0255102	<div></div>
		europe	0.0357143	<div></div>
	78to83	america	0.0561224	<div></div>
		asia	Never	
		europe	Never	
good	70to74	america	0.0102041	<div></div>
		asia	0.0306122	<div></div>
		europe	0.0459184	<div></div>
	75to77	america	0.0306122	<div></div>
		asia	0.0408163	<div></div>
		europe	0.0357143	<div></div>
	78to83	america	0.112245	<div></div>
		asia	0.0714286	<div></div>
		europe	0.0357143	<div></div>

# the Miles Per Gallon dataset

192  
Training  
Set

mpg	modelyear	maker
good	75to78	asia
bad	70to74	america
bad	75to78	europa
bad	70to74	america
bad	70to74	america
bad	70to74	asia
bad	70to74	asia

mpg	modelyear	maker		
bad	70to74	america	0.27551	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0153061	<div></div>
	75to77	america	0.153061	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0357143	<div></div>

$$\hat{P}(\text{dataset} \mid M) = \hat{P}(\mathbf{x}_1 \text{ and } \mathbf{x}_2 \dots \text{ and } \mathbf{x}_R \mid M) = \prod_{k=1}^R \hat{P}(\mathbf{x}_k \mid M)$$

$$= (\text{in this case}) = 3.4 \times 10^{-203}$$

70to74	america	0.112243	<div></div>
	asia	0.0714286	<div></div>
	europa	0.0357143	<div></div>

# Log Probabilities

Since probabilities of datasets get so small we usually use **log** probabilities

$$\log \hat{P}(\text{dataset}|M) = \log \prod_{k=1}^R \hat{P}(\mathbf{x}_k|M) = \sum_{k=1}^R \log \hat{P}(\mathbf{x}_k|M)$$

# the Miles Per Gallon dataset

192  
Training  
Set

mpg	modelyear	maker
good	75to78	asia
bad	70to74	america
bad	75to78	europa
bad	70to74	america
bad	70to74	america
bad	70to74	asia
bad	70to74	asia

mpg	modelyear	maker		
bad	70to74	america	0.27551	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0153061	<div></div>
	75to77	america	0.153061	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0357143	<div></div>

$$\log \hat{P}(\text{dataset}|M) = \log \prod_{k=1}^R \hat{P}(\mathbf{x}_k|M) = \sum_{k=1}^R \log \hat{P}(\mathbf{x}_k|M)$$

= (in this case) = -466.19

70to74	america	0.112243	<div></div>
	asia	0.0714286	<div></div>
	europa	0.0357143	<div></div>

# Summary: The Good News

- We have a way to learn a Density Estimator from data.
- Density estimators can do many good things...
  - Can sort the records by probability, and thus spot weird records (anomaly detection)
  - Can do inference:  $P(E1 | E2)$   
Automatic Doctor / Help Desk etc
  - Can perform classification, e.g.,  $p(C_k | A_1, A_2, \dots A_n)$
  - Ingredient for Bayes Classifiers (see later)

# Summary: The Bad News

- Density estimation by directly learning the joint is trivial, mindless and dangerous



# Using a test set

	Set Size	Log likelihood
Training Set	196	-466.1905
Test Set	196	-614.6157

An independent test set with 196 cars has a worse log likelihood

(actually it's a billion quintillion quintillion quintillion quintillion times less likely)

....Density estimators can overfit. And the full joint density estimator is the overfittest of them all!

# Overfitting Density Estimators

If **this** ever happens, it means there are certain combinations that we learn are impossible

mpg	modelyear	maker		
bad	70to74	america	0.27551	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0153061	<div></div>
	75to77	america	0.153061	<div></div>
		asia	0.0255102	<div></div>
		europa	0.0357143	<div></div>
	78to83	america	0.0561224	<div></div>
		asia	Never	<div></div>
		europa	Never	<div></div>
good	70to74	america	0.0102041	<div></div>
		asia	0.0255102	<div></div>

$$\log \hat{P}(\text{testset}|M) = \log \prod_{k=1}^R \hat{P}(\mathbf{x}_k|M) = \sum_{k=1}^R \log \hat{P}(\mathbf{x}_k|M)$$

$$= -\infty \text{ if for any } k \hat{P}(\mathbf{x}_k|M) = 0$$

# Using a test set

	Set Size	Log likelihood
Training Set	196	-466.1905
Test Set	196	-614.6157

The only reason that our test set didn't score -infinity is that the code is hard-wired to always predict a probability of at least one in  $10^{20}$

*We need Density Estimators that are less prone to overfitting*

# Naïve Density Estimation

The problem with the Joint Estimator is that it just mirrors the training data.

We need something which generalizes more usefully.

The naïve model generalizes strongly:

Assume that each attribute is distributed independently of any of the other attributes.

# Independently Distributed Data

- Let  $x[i]$  denote the  $i$ 'th field of record  $x$ .
- The independent distribution assumption says that for any  $i, v, u_1 u_2 \dots u_{i-1} u_{i+1} \dots u_M$

$$\begin{aligned} P(x[i] = v \mid x[1] = u_1, x[2] = u_2, \dots, x[i-1] = u_{i-1}, x[i+1] = u_{i+1}, \dots, x[M] = u_M) \\ = P(x[i] = v) \end{aligned}$$

- Or in other words,  $x[i]$  is independent of  $\{x[1], x[2], \dots, x[i-1], x[i+1], \dots, x[M]\}$
- This is often written as

$$x[i] \perp \{x[1], x[2], \dots, x[i-1], x[i+1], \dots, x[M]\}$$

# A note about independence

- Assume A and B are Boolean Random Variables.  
Then

“A and B are independent” if and only if

$$P(A|B) = P(A)$$

- “A and B are independent” is often notated as

$$A \perp B$$

# Independence Theorems

- Assume  $P(A | B) = P(A)$   
Then  
$$P(A \text{ and } B) = P(A) P(B)$$

- Assume  $P(A | B) = P(A)$   
Then  
$$P(\sim A | B) = P(\sim A)$$

- Assume  $P(A | B) = P(A)$   
Then  
$$P(B | A) = P(B)$$

- Assume  $P(A | B) = P(A)$   
Then  
$$P(A | \sim B) = P(A)$$

# Multivalued Independence

For multivalued Random Variables  $A$  and  $B$ ,

$$A \perp B$$

if and only if

$$\forall u, v : P(A = u \mid B = v) = P(A = u)$$

from which you can then prove things like...

$$\forall u, v : P(A = u \text{ and } B = v) = P(A = u)P(B = v)$$

$$\forall u, v : P(B = v \mid A = v) = P(B = v)$$



# Multivalued Independence

- Example:
  - Suppose that each record is generated by randomly shaking a green dice and a red dice
    - Dataset 1:  $A$  = red dice value,  $B$  = green dice value
    - Dataset 2:  $A$  = red dice value,  $B$  = sum of two dice values
    - Dataset 3:  $A$  = sum of the two dice values,  $B$  = difference of the two dice values
  - Which of these datasets violates the naïve assumption?

# Using the Naïve Distribution

- Once you have a Naïve Distribution you can easily compute any row of the joint distribution.
- Suppose A, B, C and D are independently distributed,

What is  $P(A \text{ and } \sim B \text{ and } C \text{ and } \sim D)$ ?

# Using the Naïve Distribution

- Once you have a Naïve Distribution you can easily compute any row of the joint distribution.
- Suppose A, B, C and D are independently distributed. What is  $P(A \text{ and } \sim B \text{ and } C \text{ and } \sim D)$ ?

$$= P(A | \sim B \wedge C \wedge \sim D) P(\sim B \wedge C \wedge \sim D)$$

$$= P(A) P(\sim B \wedge C \wedge \sim D)$$

$$= P(A) P(\sim B | C \wedge \sim D) P(C \wedge \sim D)$$

$$= P(A) P(\sim B) P(C \wedge \sim D)$$

$$= P(A) P(\sim B) P(C | \sim D) P(\sim D)$$

$$= P(A) P(\sim B) P(C) P(\sim D)$$

# Using the Naïve Distribution

- Suppose A, B, C and D are independently distributed.  
What is  $P(A \text{ and } \sim B \text{ and } C \text{ and } \sim D)$ ?

# Naïve Distribution General Case

- Suppose  $x[1], x[2], \dots, x[M]$  are independently distributed.

$$P(x[1] = u_1, x[2] = u_2, \dots, x[M] = u_M) = \prod_{k=1}^M P(x[k] = u_k)$$

- So if we have a Naïve Distribution we can construct any row of the implied Joint Distribution on demand.
- So we can do any inference
- But how do we learn a Naïve Density Estimator?

# Learning a Naïve Density Estimator

$$\hat{P}(x[i] = u) = \frac{\text{\#records in which } x[i] = u}{\text{total number of records}}$$

Another trivial learning algorithm!

# Contrast

Joint DE	Naïve DE
Can model anything	Can model only very boring distributions
No problem to model “C is a noisy copy of A”	Outside Naïve’s scope
Given 100 records and more than 6 Boolean attributes will screw up badly	Given 100 records and 10,000 multivalued attributes will be fine

# Empirical Results: “Hopeless”

The “hopeless” dataset consists of 40,000 records and 21 Boolean attributes called a,b,c, ... u. Each attribute in each record is generated 50-50 randomly as 0 or 1.

Name	Model	Parameters	LogLike
Model1	joint	submodel=gauss gausstype=general	-272625 +/- 301.109
Model2	naive	submodel=gauss gausstype=general	-58225.6 +/- 0.554747

Average test set log probability during 10 folds cross-validation

Despite the vast amount of data, “Joint” overfits hopelessly and does much worse



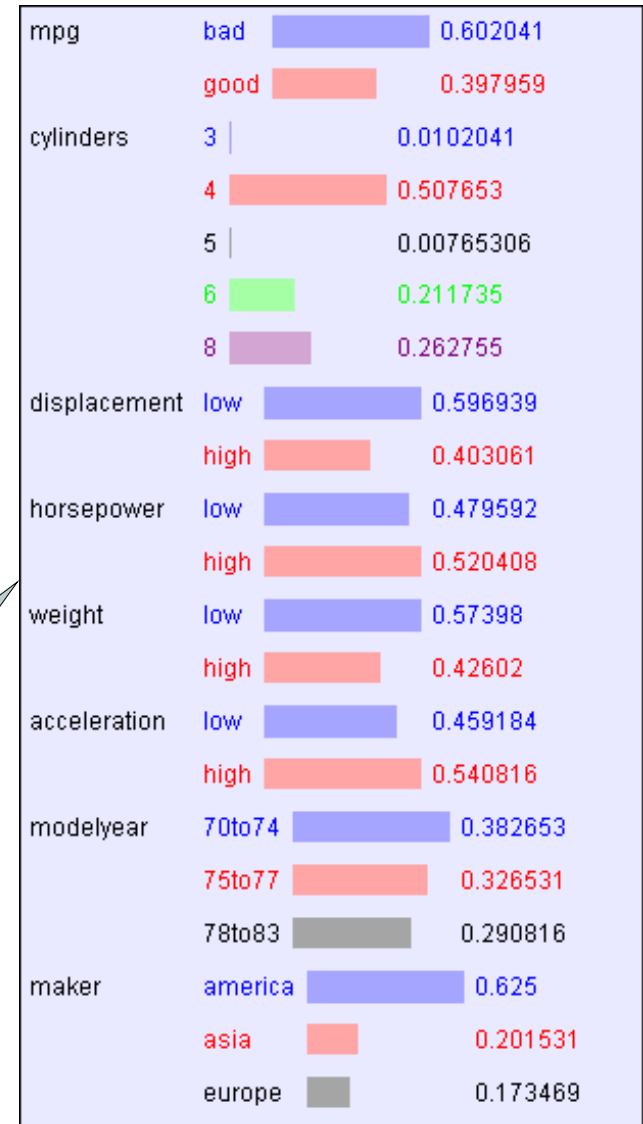
# Empirical Results: “MPG”

The “MPG” dataset consists of 392 records and 8 attributes

mpg	cylinders	displacement	horsepower	weight	acceleration	modelyear	maker
bad	3	low	low	low	low	70to74	america
							asia
							europa
						75to77	Never
						78to83	Never
					high	Never	
				high	Never		
		high	low	low		70to74	america
							asia
							europa
						75to77	america
							asia
							europa
						78to83	america
							asia
							europa
					high	Never	
				high	Never		
	high	Never					
4	low	low	low	low		70to74	america
							asia
							europa
						75to77	Never
						78to83	Never
					high	70to74	america
							asia

A tiny part of the DE learned by “Joint”

The DE learned by “Naive”

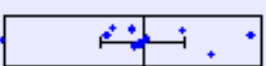



# Empirical Results: “MPG”

mpg	cylinders	displacement	horsepower	weight	acceleration	modelyear	maker
bad	3	low	low	low	low	70to74	america Never
							asia 0.00255102
							europa Never
						75to77	Never
						78to83	Never
						high	Never

A tiny part of

mpg	bad	0.602041
	good	0.397959
cylinders	3	0.0102041
	4	0.507653
	5	0.00765306

Name	Model	Parameters	LogLike	
Model1	joint	submodel=gauss gausstype=general	-472.486 +/- 77.2184	
Model2	naive	submodel=gauss gausstype=general	-257.212 +/- 3.02246	

	high	Never						
4	low	low	low	low	70to74	america	0.00255102	<div></div>
						asia	Never	
						europa	0.00255102	<div></div>
					75to77	Never		
					78to83	Never		
				high	70to74	america	0.0204082	<div></div>
						asia	0.00255102	<div></div>

The DE  
learned by  
“Naive”

acceleration	low	0.459184
	high	0.540816
modelyear	70to74	0.382653
	75to77	0.326531
	78to83	0.290816
maker	america	0.625
	asia	0.201531
	europa	0.173469

# Empirical Results: “Weight vs. MPG”

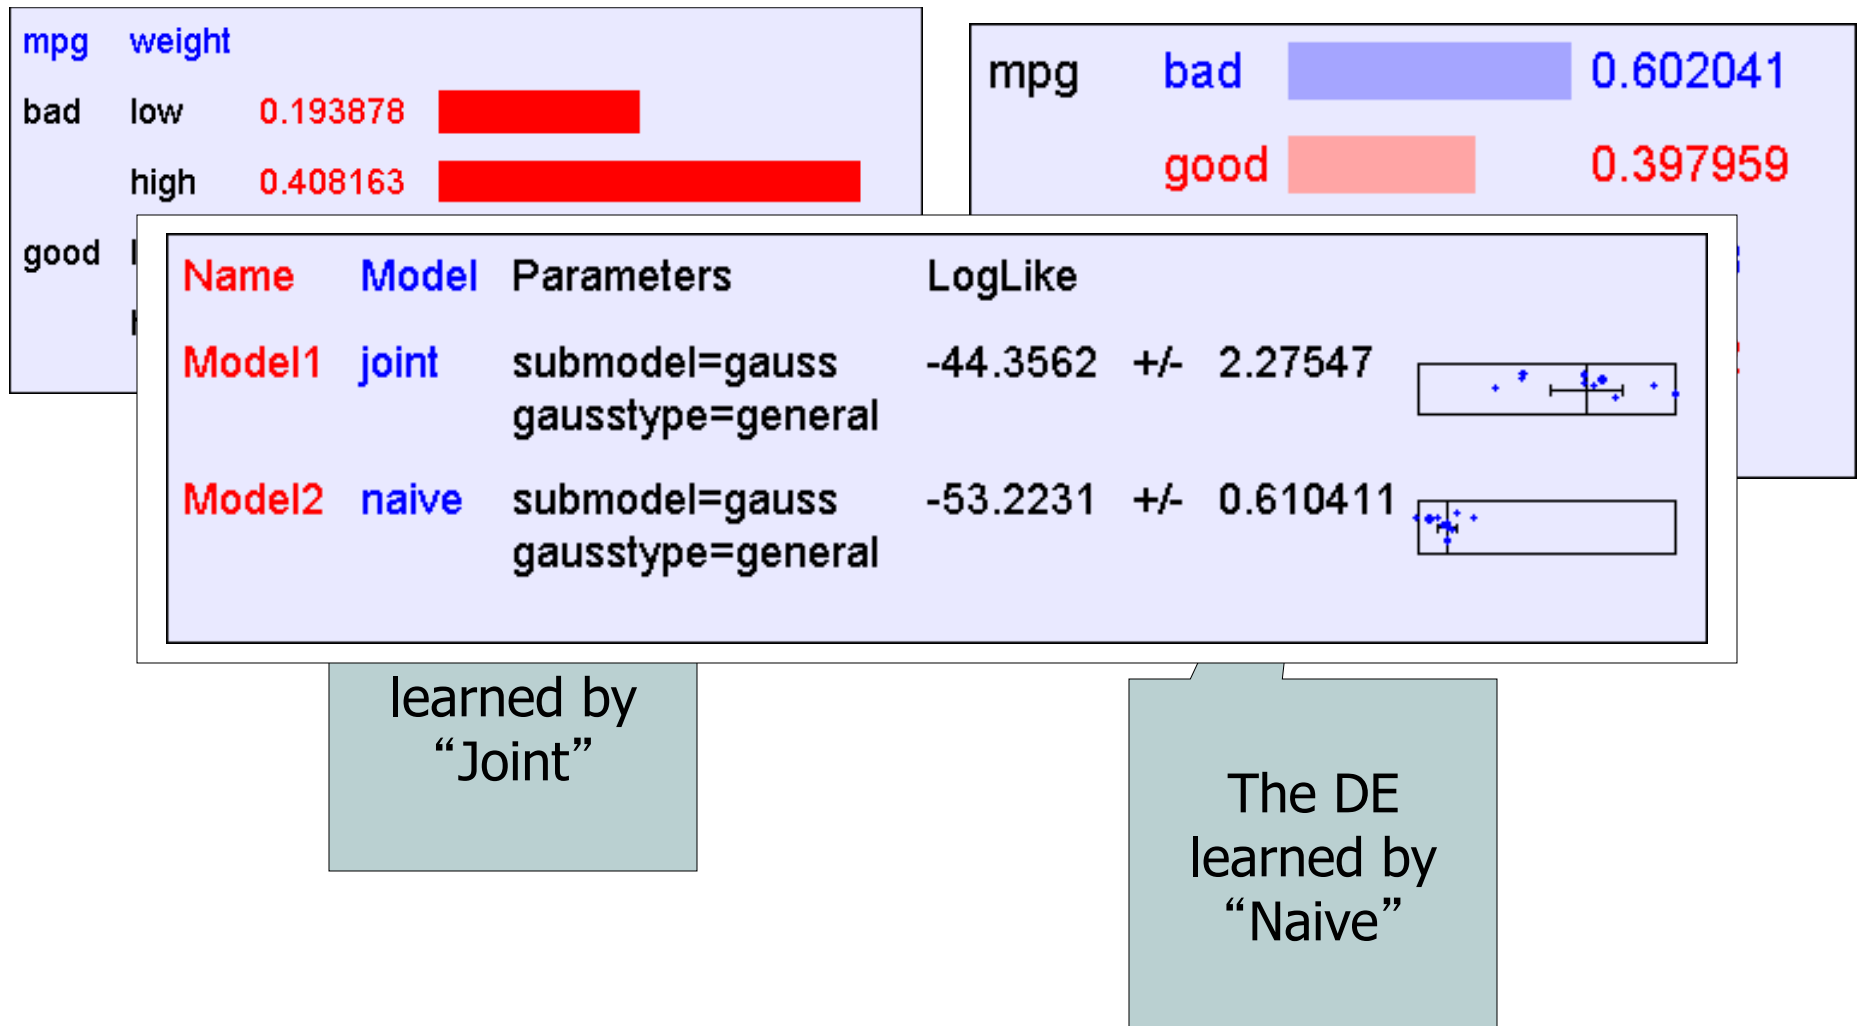
mpg	weight		
bad	low	0.193878	<div></div>
	high	0.408163	<div></div>
good	low	0.380102	<div></div>
	high	0.0178571	<div></div>

The DE  
learned by  
“Joint”

mpg	bad	<div></div>	0.602041
	good	<div></div>	0.397959
weight	low	<div></div>	0.57398
	high	<div></div>	0.42602

The DE  
learned by  
“Naive”

# Empirical Results: “Weight vs. MPG”

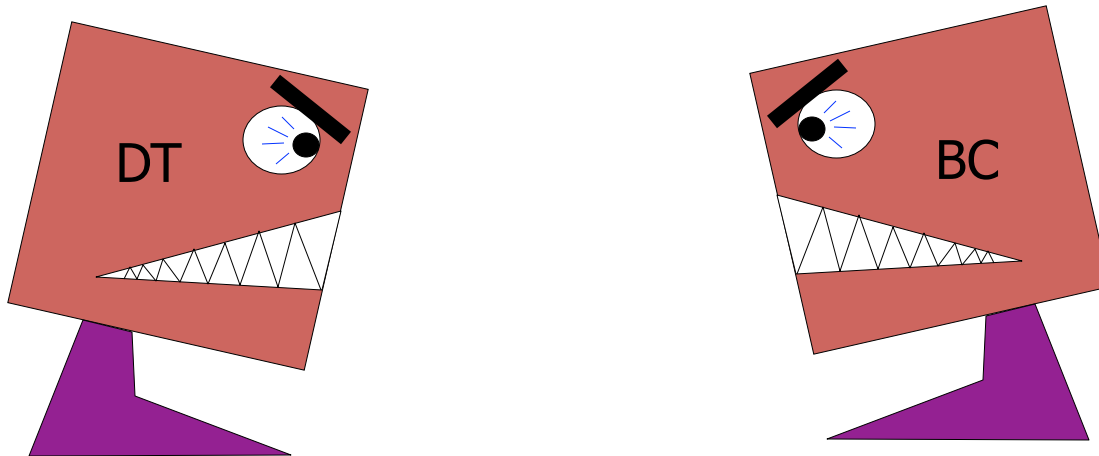


# Reminder: The Good News

- We have two ways to learn a Density Estimator from data.
- Other, vastly more impressive Density Estimators developed
  - Mixture Models, Bayesian Networks, Density Trees, Kernel Densities and many more
- Density estimators can do many good things...
  - Anomaly detection
  - Can do inference:  $P(E1|E2)$  Automatic Doctor / Help Desk etc
  - Ingredient for Bayes Classifiers

# Bayes Classifiers

- A formidable and sworn enemy of decision trees



# How to build a Bayes Classifier

- Assume you want to predict output  $Y$  which has arity  $n_Y$  and values  $v_1, v_2, \dots, v_{n_Y}$ .
- Assume there are  $m$  input attributes called  $X_1, X_2, \dots, X_m$
- Break dataset into  $n_Y$  smaller datasets called  $DS_1, DS_2, \dots, DS_{n_Y}$ .
- Define  $DS_i =$  Records in which  $Y=v_i$
- For each  $DS_i$ , learn Density Estimator  $M_i$  to model the input distribution among the  $Y=v_i$  records.

# How to build a Bayes Classifier

- Assume you want to predict output  $Y$  which has arity  $n_y$  and values  $v_1, v_2, \dots, v_{n_y}$ .
- Assume there are  $m$  input attributes called  $X_1, X_2, \dots, X_m$
- Break dataset into  $n_y$  smaller datasets called  $DS_1, DS_2, \dots, DS_{n_y}$ .
- Define  $DS_i =$  Records in which  $Y=v_i$
- For each  $DS_i$ , learn Density Estimator  $M_i$  to model the input distribution among the  $Y=v_i$  records.
- $M_i$  estimates  $P(X_1, X_2, \dots, X_m \mid Y=v_i)$



# How to build a Bayes Classifier

- Assume you want to predict output  $Y$  which has arity  $n_y$  and values  $v_1, v_2, \dots, v_{n_y}$ .
- Assume there are  $m$  input attributes called  $X_1, X_2, \dots, X_m$
- Break dataset into  $n_y$  smaller datasets called  $DS_1, DS_2, \dots, DS_{n_y}$ .
- Define  $DS_i =$  Records in which  $Y=v_i$
- For each  $DS_i$ , learn Density Estimator  $M_i$  to model the input distribution among the  $Y=v_i$  records.
- $M_i$  estimates  $P(X_1, X_2, \dots, X_m \mid Y=v_i)$
- Idea: When a new set of input values ( $X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$ ) come along to be evaluated predict the value of  $Y$  that makes  $P(X_1, X_2, \dots, X_m \mid Y=v_i)$  most likely

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1, X_2 = u_2, \dots, X_m = u_m \mid Y = v)$$

Is this a good idea?

# How to build a ~~Bayes~~ Classifier

- Assume you want to predict output  $v_{ny}$
- Assume there are  $m$  input attributes
- Break dataset into  $n_y$  smaller datasets
- Define  $DS_i$  = Records in which  $Y=v_i$
- For each  $DS_i$ , learn Density Estimator among the  $Y=v_i$  records.
- $M_i$  estimates  $P(X_1, X_2, \dots, X_m \mid Y=v_i)$

This is a Maximum Likelihood classifier.

It can get silly if some Ys are very unlikely

- Idea: When a new set of input values ( $X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$ ) come along to be evaluated predict the value of  $Y$  that makes  $P(X_1, X_2, \dots, X_m \mid Y=v_i)$  most likely

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1, X_2 = u_2, \dots, X_m = u_m \mid Y = v)$$

Is this a good idea?

# How to build a Bayes Classifier

- Assume you want to predict output  $Y$  which has  $n_Y$  possible values
  - Assume there are  $m$  input attributes called  $X_1, X_2, \dots, X_m$
  - Break dataset into  $n_Y$  smaller datasets called  $DS_i$
  - Define  $DS_i$  = Records in which  $Y=v_i$
  - For each  $DS_i$ , learn Density Estimator  $M_i$  to model the distribution of the  $X_1, X_2, \dots, X_m$  for the  $Y=v_i$  records.
  - $M_i$  estimates  $P(X_1, X_2, \dots, X_m \mid Y=v_i)$
- Much Better Idea
- Idea: When a new set of input values  $(X_1 = u_1, X_2 = u_2, \dots, X_m = u_m)$  come along to be evaluated predict the value of  $Y$  that makes  $P(Y=v_i \mid X_1, X_2, \dots, X_m)$  most likely

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v \mid X_1 = u_1, X_2 = u_2, \dots, X_m = u_m)$$

# Terminology

- MLE (Maximum Likelihood Estimator):

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)$$

- MAP (Maximum A-Posteriori Estimator):

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v \mid X_1 = u_1 \cdots X_m = u_m)$$

# Computing a posterior probability

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v \mid X_1 = u_1 \cdots X_m = u_m)$$

$$P(Y = v \mid X_1 = u_1 \cdots X_m = u_m)$$

$$\begin{aligned} &= \frac{P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)P(Y = v)}{P(X_1 = u_1 \cdots X_m = u_m)} \\ &= \frac{P(X_1 = u_1 \cdots X_m = u_m \mid Y = v)P(Y = v)}{\sum_{j=1}^{n_Y} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v_j)P(Y = v_j)} \end{aligned}$$

# Bayes Classifiers in a nutshell

1. Learn the distribution over inputs for each value  $Y$ .
2. This gives  $P(X_1, X_2, \dots, X_m \mid Y=v_i)$ .
3. Estimate  $P(Y=v_i)$  as fraction of records with  $Y=v_i$ .
4. For a new prediction:

$$\begin{aligned} Y^{\text{predict}} &= \operatorname{argmax}_v P(Y = v \mid X_1 = u_1 \cdots X_m = u_m) \\ &= \operatorname{argmax}_v P(X_1 = u_1 \cdots X_m = u_m \mid Y = v) P(Y = v) \end{aligned}$$

# Bayes Classifiers in a nutshell

- Step 1. Learn the distribution over inputs for each value  $Y$ .

- Step 2. This gives  $P(X_1, X_2, \dots$

We can use our favorite Density Estimator here.

- Step 3. Estimate  $P(Y=v_i)$  as with  $Y=v_i$ .

Right now we have two options:

- Step 4. For a new prediction

- Joint Density Estimator
- Naïve Density Estimator

$$\begin{aligned} Y^{\text{predict}} &= \operatorname{argmax}_v P(Y = v \mid X_1 = u_1 \cdots X_m = u_m) \\ &= \operatorname{argmax}_v P(X_1 = u_1 \cdots X_m = u_m \mid Y = v) P(Y = v) \end{aligned}$$

# Joint Density Bayes Classifier

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v) P(Y = v)$$

- In the case of the joint Bayes Classifier this degenerates to a very simple rule:
- $y^{\text{predict}}$  = the most common value of Y among records in which  $X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$ .
- Note that if no records have the exact set of inputs  $X_1 = u_1, X_2 = u_2, \dots, X_m = u_m$ , then  $P(X_1, X_2, \dots, X_m \mid Y = v_i) = 0$  for all values of Y.
- In that case we just have to guess Y's value



# Naïve Bayes Classifier

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v) P(Y = v)$$

- In the case of the naive Bayes Classifier this can be simplified:

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v) \prod_{j=1}^{n_Y} P(X_j = u_j \mid Y = v)$$

# An Example

Day	Outlook	Temperature	Humidity	Wind	PlayGolf
D1	sunny	hot	high	weak	no
D2	sunny	hot	high	strong	no
D3	overcast	hot	high	weak	yes
D4	rain	mild	high	weak	yes
D5	rain	cool	normal	weak	yes
D6	rain	cool	normal	strong	no
D7	overcast	cool	normal	strong	yes
D8	sunny	mild	high	weak	no
D9	sunny	cool	normal	weak	yes
D10	rain	mild	normal	weak	yes
D11	sunny	mild	normal	strong	yes
D12	overcast	mild	high	strong	yes
D13	overcast	hot	normal	weak	yes
D14	rain	mild	high	strong	no

# To Learn a Naïve Bayes Classifier from this data

Two classes:  $y=v_1$  : play golf=no  
 $y=v_2$  : play golf=yes

four attributes:

$x_1$ : three values (sunny, overcast, rain)

$x_2$ : three values (hot, mild, cool)

$x_3$ : two values (high, normal)

$x_4$ : two values (weak, strong)

# Which probabilities do we need to compute?

- $P(\text{class1} = \text{yes})$

$P(a1=\text{sunny}|y=\text{yes})$

$P(a1=\text{overcast}|y=\text{yes})$

$P(a1=\text{rain}|y=\text{yes})$

$P(a2=\text{hot}|y=\text{yes})$

$P(a2=\text{mild}|y=\text{yes})$

$P(a2=\text{cool}|y=\text{yes})$

$P(a3=\text{high}|y=\text{yes})$

$P(a3=\text{normal}|y=\text{yes})$

$P(a4=\text{weak}|y=\text{yes})$

$P(a4=\text{strong}|y=\text{yes})$

$P(\text{class2}=\text{no})$

$P(a1=\text{sunny}|y=\text{no})$

$P(a1=\text{overcast}|y=\text{no})$

$P(a1=\text{rain}|y=\text{no})$

$P(a2=\text{hot}|y=\text{no})$

$P(a2=\text{mild}|y=\text{no})$

$P(a2=\text{cool}|y=\text{no})$

$P(a3=\text{high}|y=\text{no})$

$P(a3=\text{normal}|y=\text{no})$

$P(a4=\text{weak}|y=\text{no})$

$P(a4=\text{strong}|y=\text{no})$

# Reorder according to class label

Day	Outlook	Temperature	Humidity	Wind	Play Golf
D1	sunny	hot	high	weak	no
D2	sunny	hot	high	strong	no
D3	overcast	hot	high	weak	yes
D4	rain	mild	high	weak	yes
D5	rain	cool	normal	weak	yes
D6	rain	cool	normal	strong	no
D7	overcast	cool	normal	strong	yes
D8	sunny	mild	high	weak	no
D9	sunny	cool	normal	weak	yes
D10	rain	mild	normal	weak	yes
D11	sunny	mild	normal	strong	yes
D12	overcast	mild	high	strong	yes
D13	overcast	hot	normal	weak	yes
D14	rain	mild	high	strong	no

# Classification Step

Given a new case/object:

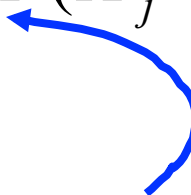
outlook=sunny,  
temperature=cool,  
humid=high,  
wind = strong

Question: whether to play or not to play golf?

# Naïve Bayes Classifier

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(X_1 = u_1 \cdots X_m = u_m \mid Y = v) P(Y = v)$$

- In the case of the naive Bayes Classifier this can be simplified:

$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} P(Y = v) \prod_{j=1}^{n_Y} P(X_j = u_j \mid Y = v)$$


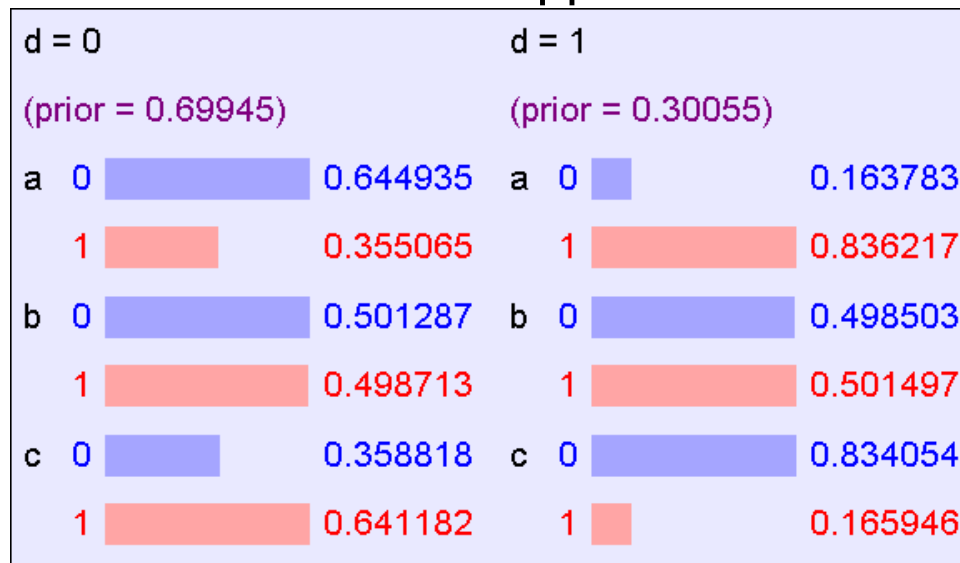
Technical Hint:

If you have 10,000 input attributes **that** product will underflow in floating point math. You should use logs:

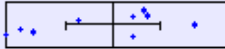
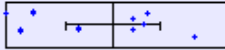
$$Y^{\text{predict}} = \underset{v}{\operatorname{argmax}} \left( \log P(Y = v) + \sum_{j=1}^{n_Y} \log P(X_j = u_j \mid Y = v) \right)$$

# Naive BC Results: “Logical”

The “logical” dataset consists of 40,000 records and 4 Boolean attributes called a,b,c,d where a,b,c are generated 50-50 randomly as 0 or 1.  $d = a \text{ and } \sim c$ , except that in 10% of records it is flipped

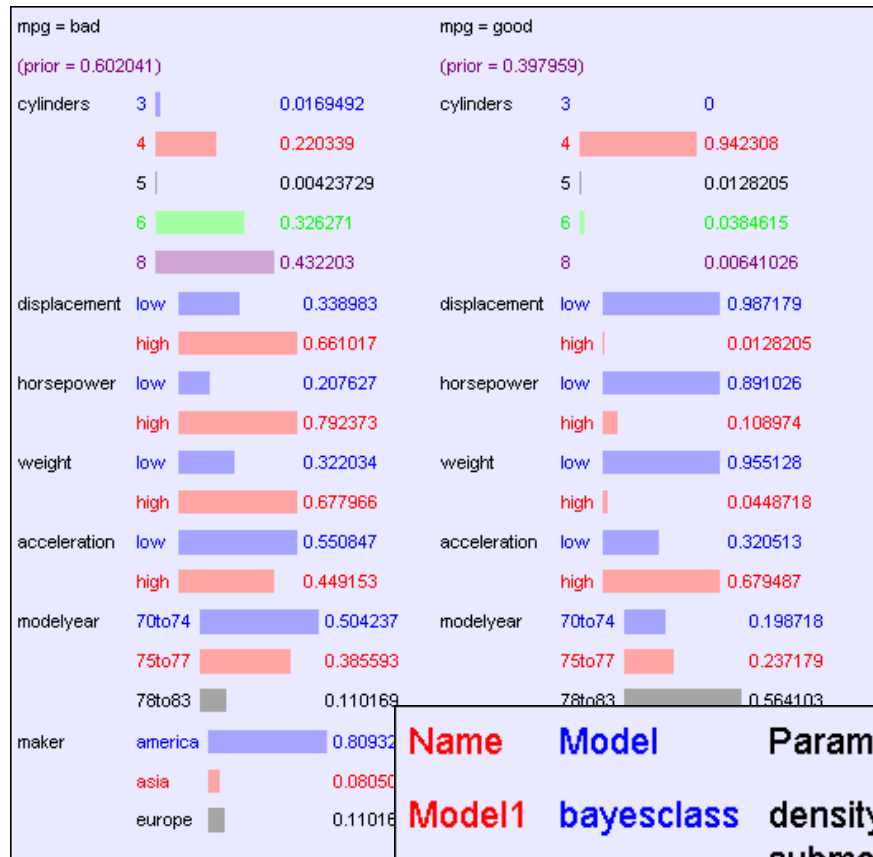


The Classifier  
learned by  
“Naive BC”

Name	Model	Parameters	FracRight	
Model1	bayesclass	density=joint submodel=gauss gausstype=general	0.90065 +/- 0.00301897	
Model2	bayesclass	density=naive submodel=gauss gausstype=general	0.90065 +/- 0.00301897	





# BC Results: “MPG”: 392 records



The Classifier learned by “Naive BC”

Name	Model	Parameters	FracRight	
Model1	bayesclass	density=joint submodel=gauss gausstype=general	0.885256 +/- 0.0247796	
Model2	bayesclass	density=naive submodel=gauss gausstype=general	0.852372 +/- 0.0400495	

# BC Results: “MPG”: 40 records

Name	Model	Parameters	FracRight	
Model1	bayesclass	density=joint submodel=gauss gausstype=general	0.725	+/- 0.114333 
Model2	bayesclass	density=naive submodel=gauss gausstype=general	0.8	+/- 0.122227 

# Classify text with naïve Bayes classifier

- Why?
  - Learn which news articles are of interest
  - Learn to classify web pages by topic
  - Spam control...
- Naïve Bayes is among the most effective algorithms

What attributes shall we use to represent text documents?

# Text Classification – data formulation

- Class label:  
Target concept Interesting?  
Document → {class1=yes, class2=no}
- represent each document by vector of words (one attribute per word position in document)
  - Remove stopwords, numbers, tags, single letters, ...
  - Change all words to lower case
  - Stemming (only retain roots)
  - Remove words appeared only once

# Naïve Bayes Classifier for Text Classification

- Build classifier: estimate  
 $P(\text{class1=yes}), P(\text{class2=no}),$   
 $P(\text{doc} | \text{class1=yes}), P(\text{doc} | \text{class2=no})$

conditional independence assumption:

$$P(\text{doc} | \text{class}_j) = \prod_{i=1}^{\text{length}(\text{doc})} P(a_i = w_k | \text{class}_j)$$

Probability word  
in position  $i$  is  $w_k$   
for class <sub>$j$</sub>

# Naïve Bayes Classifier for Text Classification

- Additional assumption: **positional independence assumption**

drop word positioning

$$P(a_i = w_k | class_j) = P(a_m = w_k | class_j), \text{ for all } i, m$$

Therefore,

$$\begin{aligned} P(doc | class_j) &= \prod_{i=1}^{length(doc)} P(a_i = w_k | class_j) \\ &= \prod_i^{length(doc)} P(w_i | class_j) \end{aligned}$$

# Steps in Learning Naïve Bayes Text Classifier

- Collect all words and other tokens that occur in examples
- Vocabulary = all distinct words and other tokens in the examples
- Calculate  $P(\text{class}_j)$  and  $P(w_k | \text{class}_j)$  for each target value  $\text{class}_j$  :
  - $\text{doc}_j$  = subset of document examples for which the target value is  $\text{class}_j$
  - $P(\text{class}_j) = |\text{doc}_j| / |\text{all document examples}|$
  - $\text{text}_j \leftarrow$  a single document created by concatenating all members of  $\text{doc}_j$

# Steps in Learning Naïve Bayes Text Classifier

- $n$  = total number of words in  $\text{Text}_j$  (counting duplicate words multiple times)
- for each word  $w_k$  in Vocabulary

$n_k$  = number of times word  $w_k$  occurs in  $\text{Text}_j$

$$P(w_k | \text{class}_j) = \frac{(n_k + 1)}{n + |\text{vocabulary}|}$$



# Steps in Classifying a Document using the Naïve Test Classifier

- Positions = all word positions in the document that contain tokens found in **Vocabulary**
- Return  $v_{NB}$ , where

$$v_{NB} = \operatorname{argmax}_j P(class_j) \prod_{i \in positions} P(w_i | class_j)$$

# Example Application: Classify newsgroup documents

Given 1000 training documents from each group,  
learn to classify new documents according to  
which newsgroup it came from:

comp.graphics

misc.forsale

comp.os.ms-windows.misc

rec.autos

comp.sys.ibm.pc.hardware

rec.motorcycles

comp.sys.mac.hardware

rec.sport.hockey

....

Result: Naïve Bayes obtained 89% classification accuracy