

Data Mining



Logistic Regression

Notes adapted from Logistic Regression in Python Course

Classification Problem

- Supervised vs. Unsupervised Classification
- Supervised classification
 - Regression: predict the value of a function
 - Classification: put data into classes
 - Logistic regression

MNIST






Image classification



Linear Regression

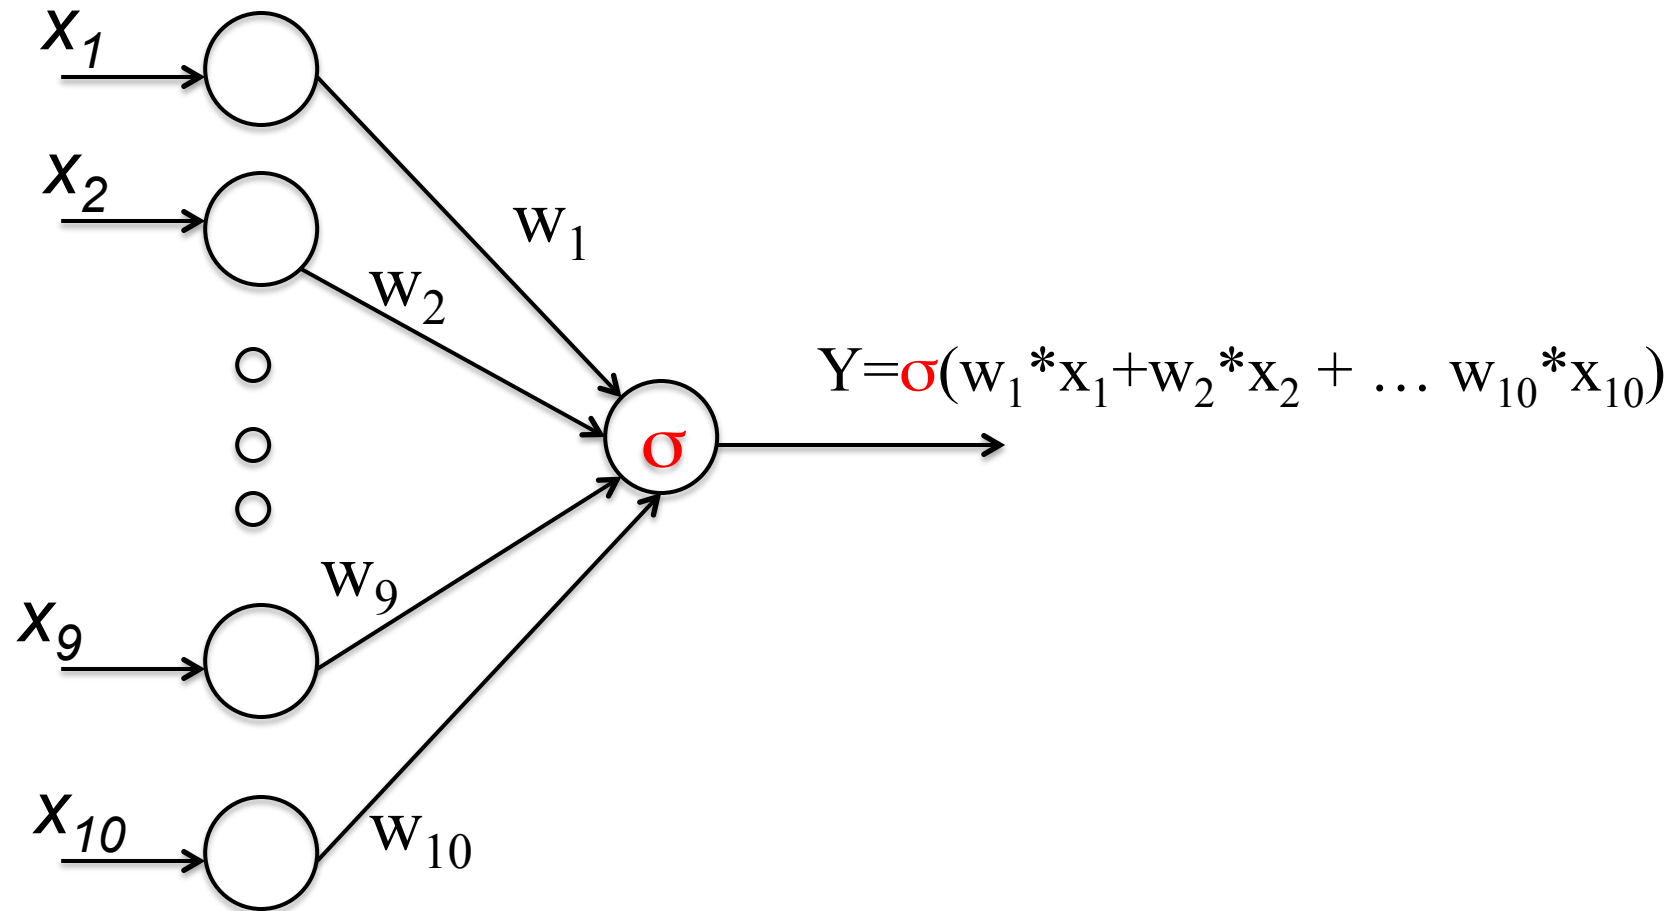
$y=mx+b$, or
 $ax+by+c=0$
 $a=1, b=-1, c=0$
 $x-y=0$



Regression


- $h(x) = w_0 + w_1x_1 + w_2x_2$, $h()$ is a linear combination of the components of x
 - In vector form : $h(x) = w^T x$
- The class separating function:
 - In 2-dimensions: a line
 - In 3-dimensions: a plane
 - In >3 dimension: hyperplane

Logistic Regression



10 features

Logistic Regression



$$y = \sum_{i=0}^N w_i * x_i$$


$$y = \sigma(w^T x)$$

10 features

Activation Functions

- Tanh()
$$f(x) = \frac{2}{1 + e^{-2x}} - 1$$
$$f'(x) = 1 - f(x)^2$$
- Sigmoid/Logistic
$$f(x) = \frac{1}{1 + e^{(-x)}}$$
$$f'(x) = f(x)[1 - f(x)]$$
- Bipolar Sigmoid
$$f(x) = \frac{2}{1 + e^{(-x)}} - 1$$
$$f'(x) = \frac{1}{2}[1 + f(x)][1 - f(x)]$$

Sigmoid Function for Classification



$$y = \sigma(w^T x)$$

if $\sigma(w^T x) > 0.5$, predict class 1

else if $\sigma(w^T x) < 0.5$, predict class 0

Demo logistic1.py

Real Data Application

- Features:
 - Is_mobile (0/1)
 - N_Products_viewed (int >0)
 - Visit_duration (real ≥ 0)
 - Is_returning_visitor (0/1)
 - Time_of_day (1, 2, 3, 4 = 4 24 hour split into 4 categories)
 - Class: User_action (bounce/add_to_cart/begin_checkout)

Demo ecommerce_data.csv

Data Preprocessing

- Data pre-processing
- One-hot encoding:

12am-6am , 6am-12pm, 12pm-6pm, 6pm - 12am

1	0	0	0
0	1	0	0
0	0	1	0
0	0	0	1

- Use Z-score for numeric features:
n_products_viewed and visitor_duration

Demo process.py

Prediction on Ecommerce Data

- Random weights are used. How good is the prediction accuracy?
- Two things to come:
 - How to evaluate the quality of the logistic regression model learned?
 - How to improve the prediction accuracy?

Demo `logistic_predict.py`

Cross Entropy Error Cost Function

- Logistic Regression Error
 - 0 if correct, >0 if not correct, more wrong ==> bigger cost

- Cross-Entropy Error cost function

$$J = -(t \log(y) + (1 - t) \log(1 - y))$$

t is the target, y is the predicted value

t=1, y=1 --> 0

t=0, y=0 --> 0

t=1, y=0.9 --> 0.11

t=1, y=0.5 --> 0.69

t=1, y=0.1 --> 2.3

Demo logistic2.py

Minimizing the Cost Function

Multiple Training Examples

$$J = - \sum_{n=1}^N t_n \log(y_n) + (1 - t_n) \log(1 - y_n)$$

- To minimizing the cost function over the entire data set

$$\frac{\partial J}{\partial w_i} = 0$$

- Generally, there is no closed form solution for this minimization problem, except for special cases

Minimizing the Cost Function

Gradient Descent

Idea: take small steps in direction of derivative

Step size == learning rate (1 for this example)

Ex:

$$w = -2$$

$$w = -2 - 1*(-1) = -1$$

Now we're closer to the optimal point! ($w=0$)

Note: slope is 0 at the bottom, so no more changes will occur



Minimizing the Cost Function

Gradient Descent for Logistic Regression

$$J = - \sum_{n=1}^N t_n \log(y_n) + (1 - t_n) \log(1 - y_n)$$

Split into 3 derivatives:

$$\frac{\partial J}{\partial w_i} = \sum_{n=1}^N \frac{\partial J}{\partial y_n} \frac{\partial y_n}{\partial a_n} \frac{\partial a_n}{\partial w_i}$$
$$a_n = w^T x_n$$

Minimizing the Cost Function

Derivatives

$$J = - \sum_{n=1}^N t_n \log(y_n) + (1 - t_n) \log(1 - y_n)$$

$$\frac{\partial J}{\partial y_n} = - t_n \frac{1}{y_n} + (1 - t_n) \frac{1}{1 - y_n} (-1)$$

Minimizing the Cost Function

Derivatives

$$y_n = \sigma(a_n) = \frac{1}{1+e^{-a_n}}$$

$$\frac{\partial y_n}{\partial a_n} = \frac{-1}{(1+e^{-a_n})^2} (e^{-a_n})(-1)$$

$$\frac{\partial y_n}{\partial a_n} = \frac{e^{-a_n}}{(1+e^{-a_n})^2} = \frac{1}{1+e^{-a_n}} \frac{e^{-a_n}}{1+e^{-a_n}} = y_n(1 - y_n)$$

Minimizing the Cost Function

Derivatives

$$a_n = w^T x_n$$

$$a_n = w_0 x_{n0} + w_1 x_{n1} + w_2 x_{n2} + \dots$$

$$\frac{\partial a_n}{\partial w_i} = x_{ni}$$

Minimizing the Cost Function

Putting them all together

$$\frac{\partial J}{\partial w_i} = - \sum_{n=1}^N \frac{t_n}{y_n} y_n (1 - y_n) x_{ni} - \frac{1-t_n}{1-y_n} y_n (1 - y_n) x_{ni}$$

$$\frac{\partial J}{\partial w_i} = - \sum_{n=1}^N t_n (1 - y_n) x_{ni} - (1 - t_n) y_n x_{ni}$$

$$\frac{\partial J}{\partial w_i} = - \sum_{n=1}^N [t_n - t_n y_n - y_n + t_n y_n] x_{ni}$$

$$\frac{\partial J}{\partial w_i} = \sum_{n=1}^N (y_n - t_n) x_{ni}$$

Weight Updates

Repeat for M epoch:
on each epoch:

$$w_i += \nabla w,$$

$$\nabla w = \lambda * (y - t) * x_i$$

λ : Learning Rate \rightarrow step size

t : target class value

y: predicted class value

Demo logistic3.py
Demo logistic_train.py