




# Clustering

---

## Part Two: Hierarchical clustering approaches

# Hierarchical Clustering

- Use distance matrix as clustering criteria. This method does not require the number of clusters  $k$  as an input, but needs a termination condition





# Single-link vs. Complete-link

---

- difference: the way to characterize the similarity between a pair of clusters
  - **single link:** *minimum* of the distances between all pairs of patterns drawn from the two clusters
  - **complete link:** *maximum* of the distances between all pairs of patterns drawn from the two clusters
  - **Complete link:** average of the distances between all pairs off
- Both use agglomerative clustering control structure



# Agglomerative clustering

---

- Step 1: place each pattern in its own cluster  
construct a list of inter-pattern distances for all distinct unordered pairs of patterns, and sort this list in ascending order
- Step through the sorted list of distances, forming for each distinct dissimilarity value  $d_k$ , a graph on the patterns where pairs of patterns closer than  $d_k$  are connected by a graph edge.  
If all patterns are members of a completely connected graph, stop.




# ***Dendrogram***

---

## **A *Dendrogram* Shows How the Clusters are Merged Hierarchically:**

- Decompose data objects into a several levels of nested partitioning (tree of clusters), called a dendrogram.
- A clustering of the data objects is obtained by cutting the dendrogram at the desired level, then each connected component forms a cluster.





# Practice question




---

- Cluster the following six objects, using single-link and complete link agglomerative clustering methods:


	Gender	Age	time	Fever	cough
Obj1:	F	23	2	Y	N
Obj2:	M	2	0.5	N	N
Obj3:	F	15	3	Y	Y
Obj4:	F	18	0.5	Y	N
Obj5:	M	58	4	N	Y
Obj6:	F	44	14	N	Y

# AGNES (Agglomerative Nesting)

- Implemented in statistical analysis packages, e.g., Splus
- Use the Single-Link method and the dissimilarity matrix.
- Merge nodes that have the least dissimilarity
- Go on in a non-descending fashion
- Eventually all nodes belong to the same cluster








# DIANA (Divisive Analysis)

- Inverse order of AGNES
- Eventually each node forms a cluster on its own





# More on Hierarchical Clustering Methods

---

- **Major weakness of agglomerative clustering methods**
  - do not scale well: time complexity of at least  $O(n^2)$ , where  $n$  is the number of total objects
  - can never undo what was done previously
- **Integration of hierarchical with distance-based clustering**
  - BIRCH (1996): uses CF-tree and incrementally adjusts the quality of sub-clusters
  - CURE (1998): selects well-scattered points from the cluster and then shrinks them towards the center of the cluster by a specified fraction
  - ...



# Introduction to BIRCH

---

- Only works with "metric" attributes
  - *Must have Euclidean coordinates*
- Designed for very large data sets
  - *Time and memory constraints are explicit*
  - *Treats dense regions of data points as sub-clusters*
    - *Not all data points are important for clustering*
  - *Only one scan of data is necessary*



# Introduction to BIRCH

---

- Incremental, distance-based approach
  - *Does not need the whole data set in advance*
  - *Unique approach: distance based algorithms generally need all the data points to work*
- Does not assume that the probability distributions on attributes is independent



# Balanced Iterative Reducing and Clustering using Hierarchies (BIRCH)

---

- Incrementally construct a CF (Clustering Feature) tree, a hierarchical data structure for multiphase clustering
  - Phase 1: scan DB to build an initial in-memory CF tree (a multi-level compression of the data that tries to preserve the inherent clustering structure of the data)
  - Phase 2: use an arbitrary clustering algorithm to cluster the leaf nodes of the CF-tree
- *Scales linearly*: finds a good clustering with a single scan and improves the quality with a few additional scans
- *Weakness*: handles only numeric data, and sensitive to the order of the data record.


# Clustering Feature Vector

**Clustering Feature:**  $CF = (N, \overrightarrow{LS}, SS)$

**$N$ :** Number of data points

$$LS: \sum_{i=1}^N \overrightarrow{X_i}$$

$$SS: \sum_{i=1}^N \overrightarrow{X_i}^2$$



$$CF = (5, (16,30), (54,190))$$

(3,4)

(2,6)

(4,5)

(4,7)

(3,8)



## The algorithm: background

---

*Given a cluster of instances  $\{\vec{X}_i\}$ , we define the **centroid**, the **radius**, and the **diameter**:*

$$\vec{X}_0 = \frac{\sum_{i=1}^N \vec{X}_i}{N}$$

$$R = \left( \frac{\sum_{i=1}^N (\vec{X}_i - \vec{X}_0)^2}{N} \right)^{\frac{1}{2}}$$

$$D = \left( \frac{\sum_{i=1}^N \sum_{j=1}^N (\vec{X}_i - \vec{X}_j)^2}{N(N-1)} \right)^{\frac{1}{2}}$$





## The algorithm: background

---

*We define the **Euclidean** and **Manhattan** distance between any two clusters as:*

$$D0 = ((\vec{X}0_1 - \vec{X}0_2)^2)^{\frac{1}{2}}$$

$$D1 = |\vec{X}0_1 - \vec{X}0_2| = \sum_{i=1}^d |\vec{X}0_1^{(i)} - \vec{X}0_2^{(i)}|$$



# The algorithm: background

*We define the average inter-cluster, the average intra-cluster, and the variance increase distances as:*

$$D2 = \left( \frac{\sum_{i=1}^{N_1} \sum_{j=N_1+1}^{N_1+N_2} (\vec{X}_i - \vec{X}_j)^2}{N_1 N_2} \right)^{\frac{1}{2}}$$

$$D3 = \left( \frac{\sum_{i=1}^{N_1+N_2} \sum_{j=1}^{N_1+N_2} (\vec{X}_i - \vec{X}_j)^2}{(N_1 + N_2)(N_1 + N_2 - 1)} \right)^{\frac{1}{2}}$$

$$D4 = \left( \sum_{k=1}^{N_1+N_2} \left( \vec{X}_k - \frac{\sum_{l=1}^{N_1+N_2} \vec{X}_l}{N_1+N_2} \right)^2 - \sum_{i=1}^{N_1} \left( \vec{X}_i - \frac{\sum_{l=1}^{N_1} \vec{X}_l}{N_1} \right)^2 - \sum_{j=N_1+1}^{N_1+N_2} \left( \vec{X}_j - \frac{\sum_{l=N_1+1}^{N_1+N_2} \vec{X}_l}{N_2} \right)^2 \right)^{\frac{1}{2}}$$



# The algorithm: background

---

- Preprocessing data is optional
- Can not affect relative placement
  - *If point A is left of B, then after preprocessing A must still be to the left of B*
- Avoids bias caused by dimensions with a large spread
- Large spread may naturally describe data



## The algorithm: CF

---

*A **Clustering Feature** (CF) summarizes a sub-cluster of data points.*

Given a cluster  $\{\vec{X}_1, \vec{X}_2, \dots, \vec{X}_N\}$

$$\mathbf{CF} = (N, \vec{LS}, SS)$$

$N$  is the number of data points

$$\vec{LS} = \sum_{i=1}^N \vec{X}_i$$

$$SS = \sum_{i=1}^N \vec{X}_i^2$$

$$\mathbf{CF}_1 + \mathbf{CF}_2 = (N_1 + N_2, \vec{LS}_1 + \vec{LS}_2, SS_1 + SS_2)$$

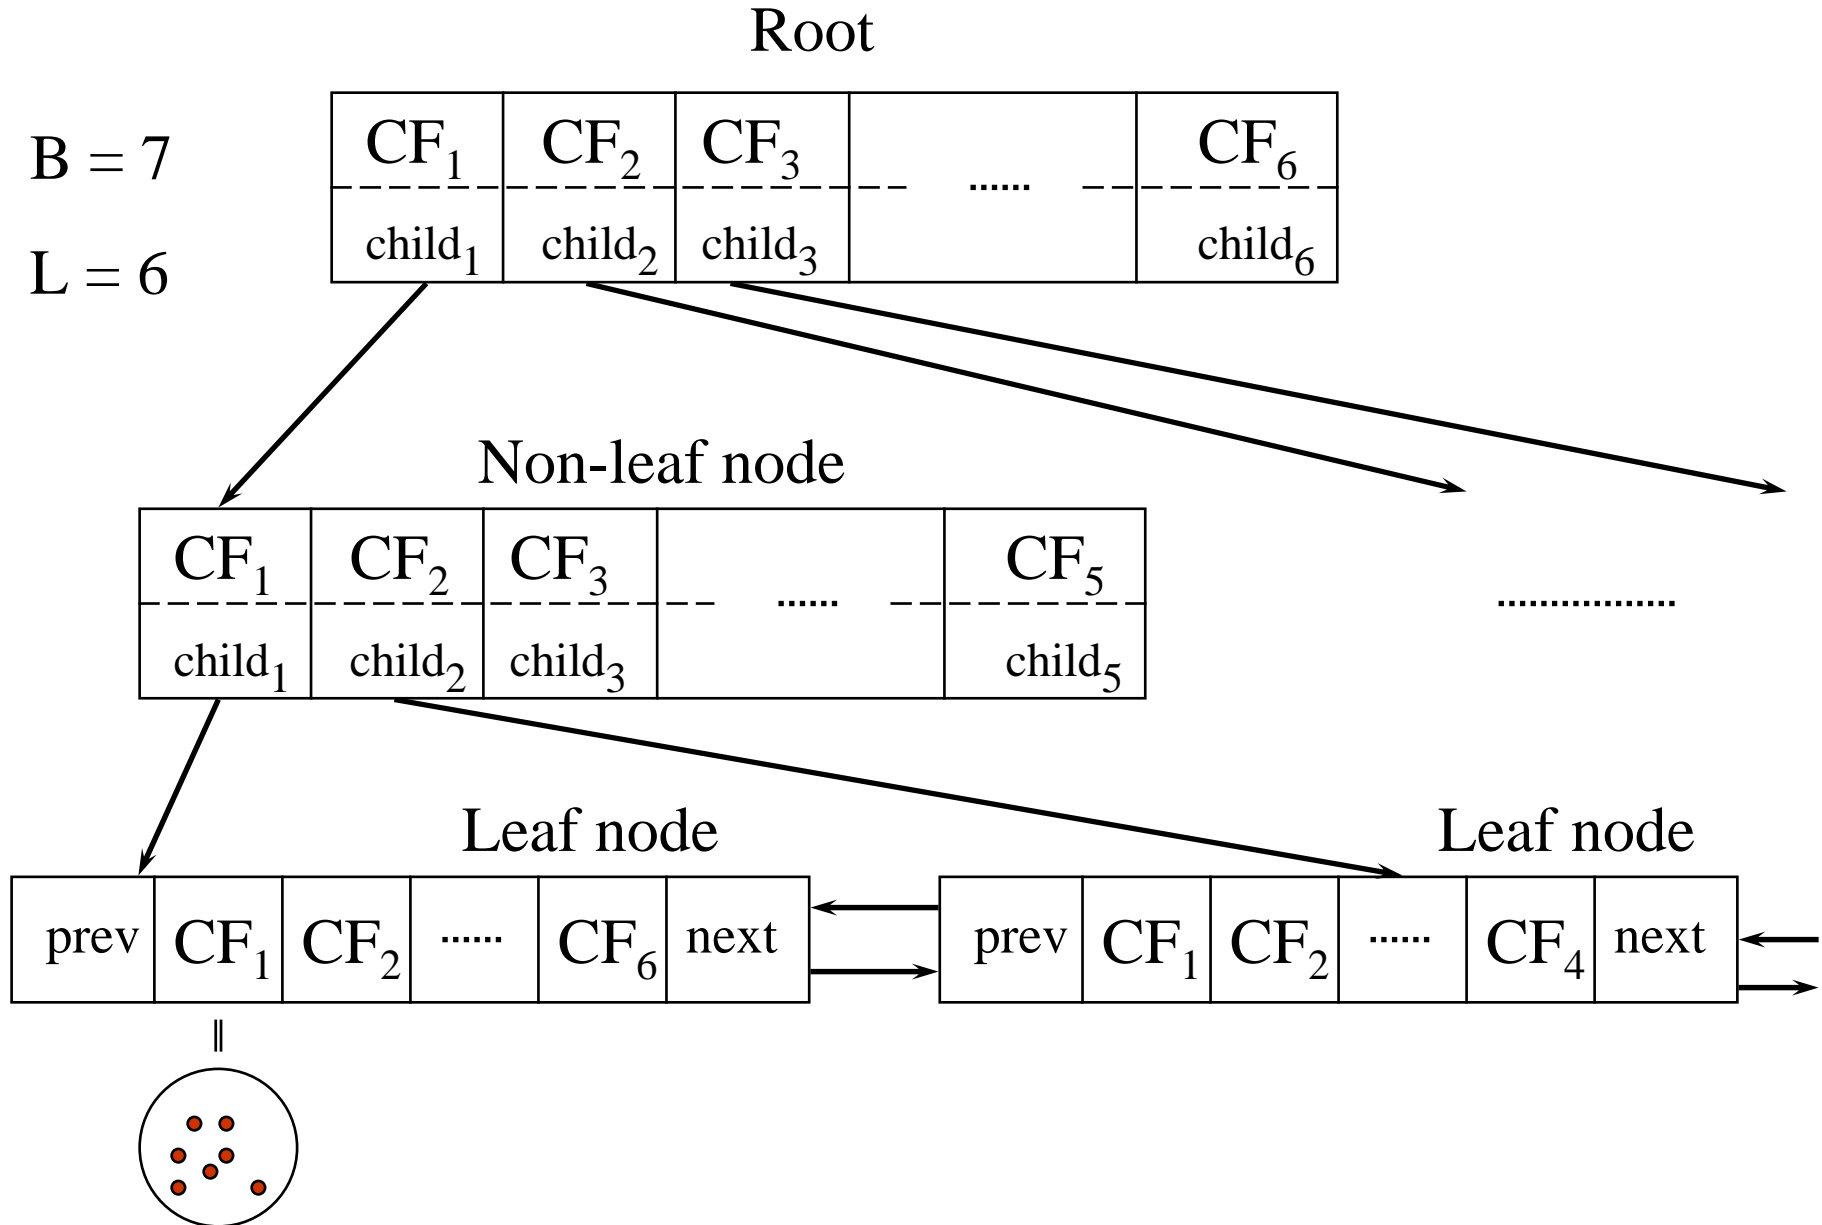


## The algorithm: CF

---

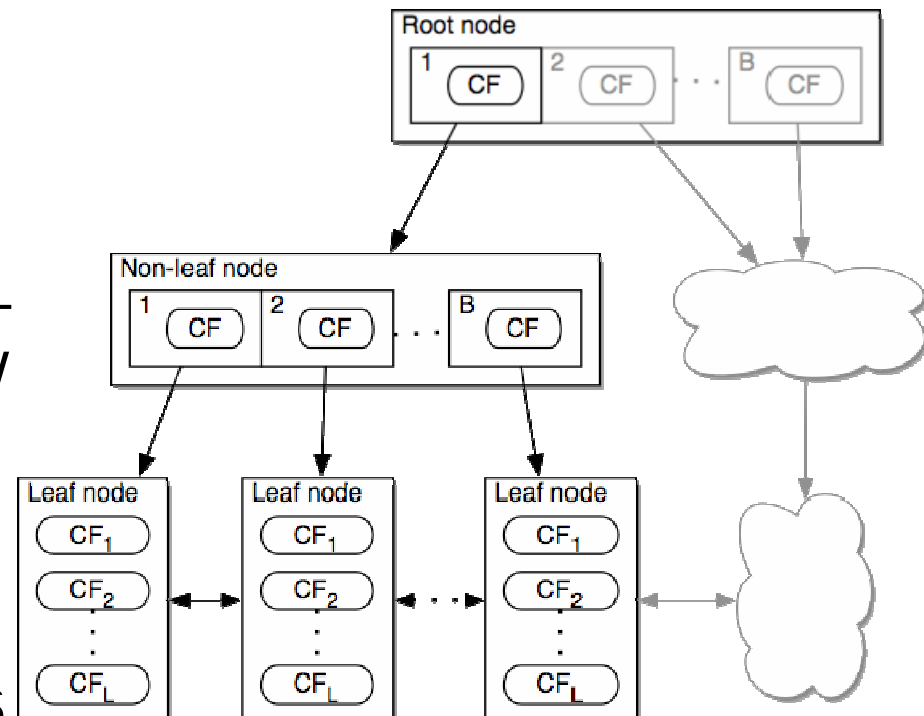
- CF entry is more compact
  - *Stores significantly less than all of the data points in the sub-cluster*
- A CF entry has enough information to calculate D0-D4
- Additivity theorem allows us to incrementally merge sub-clusters

# The algorithm: CF-tree



# The algorithm: CF-tree

- Each non-leaf node has at most  $B$  entries
- Each leaf node has at most  $L$  CF entries which each satisfy threshold  $T$
- Node size is determined by dimensionality of data points and input parameter  $P$  (page size)





# The algorithm: CF-tree insert

---

- Recurse down from root
  - *Choose the "closest" CF and go to that node*
- Modify the leaf
  - *If the closest CF in the leaf can not absorb, make a new CF entry. If there is no room, split the node*
- Traverse back up
  - *Modifying CFs or splitting nodes*





# The algorithm: CF-tree rebuild

---

- If we run out of space, increase  $T$ 
  - *By increasing the threshold, CFs absorb more data*
- Rebuilding "pushes" CFs over
  - *The larger  $T$  allows different CFs to group together*
- Reducibility theorem
  - *Increasing  $T$  will result in a CF-tree as small or smaller than the original*



# The algorithm: BIRCH

---

- Phase 1: Load data into memory
  - *Build a CF-tree with the data*
- Phase 2: Condense data
  - *Rebuild the CF-tree with a larger T*
  - *Condensing is optional*



# The algorithm: BIRCH

---

- Phase 3: Global clustering
  - *Use existing clustering algorithm on CF entries*
  - *Helps fix problem where natural clusters span nodes*
- Phase 4: Cluster refining
  - *Do additional passes over the data set and reassign data points to the closest centroid from phase 3*
  - *Refining is optional*



# The algorithm: BIRCH

---

- Why have optional phases?
  - *Phase 2 allows us to resize the data set so Phase 3 runs on an optimally sized data set*
  - *Phase 4 fixes a problem with CF-trees where some data points may be assigned to different leaf entries*
  - *Phase 4 will always converge to a minimum*
  - *Phase 4 allows us to discard outliers*



# The algorithm: BIRCH

---

- Rebuild CF-tree with smallest  $T$ 
  - *Start with  $T=0$  and try rebuilding the tree*
- Get rid of outliers
  - *Write outliers to special place outside of the tree*
- Delayed split
  - *Treat data points that force a split like outliers*



# Experimental results

---

- Input parameters:

- *Memory ( $M$ ): 5% of data set*
- *Disk space ( $R$ ): 20% of  $M$*
- *Distance equation:  $D2$*
- *Quality equation: weighted average diameter ( $D$ )*
- *Initial threshold ( $T$ ): 0.0*
- *Page size ( $P$ ): 1024 bytes*



# Experimental results

---

- The Phase 3 algorithm used is called an agglomerative Hierarchical Clustering (HC) algorithm
- One refinement pass
  - *Outlier discarding off*
- Delay-split is on
  - *This is what we use disk space  $R$  for*



# Experimental results

---


- Create 3 synthetic data sets for testing
  - *Also create an ordered copy for testing input order*
- KMEANS and CLARANS require entire data set to be in memory
  - *Initial scan is from disk, subsequent scans are in memory*





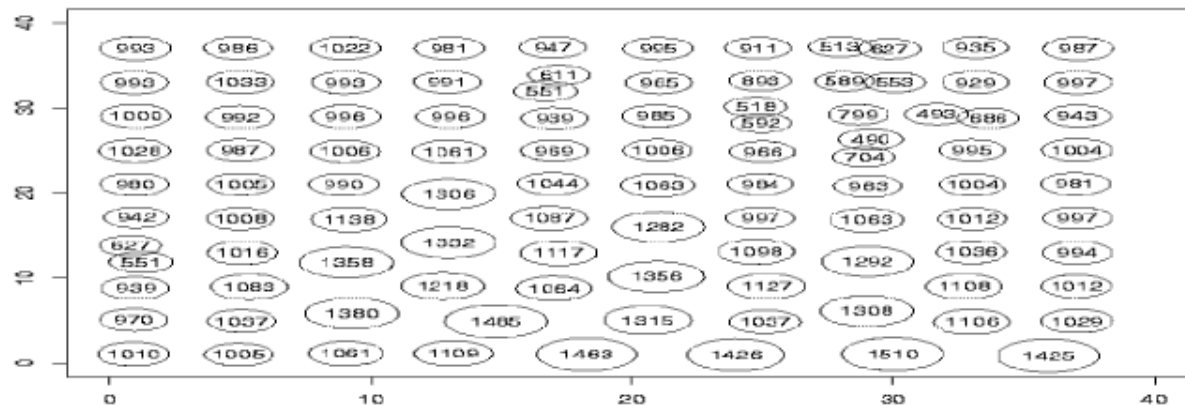
# Experimental results

*Intended clustering*



# Experimental results


## *KMEANS clustering*



<i>DS</i>	<i>Time</i>	<i>D</i>	<i># Scan</i>	<i>DS</i>	<i>Time</i>	<i>D</i>	<i># Scan</i>
1	43.9	2.09	289	10	33.8	1.97	197
2	13.2	4.43	51	20	12.7	4.20	29
3	32.9	3.66	187	30	36.0	4.35	241

# Experimental results


## *CLARANS clustering*



<i>DS</i>	<i>Time</i>	<i>D</i>	<i># Scan</i>	<i>DS</i>	<i>Time</i>	<i>D</i>	<i># Scan</i>
1	932	2.10	3307	10	794	2.11	2854
2	758	2.63	2661	20	816	2.31	2933
3	835	3.39	2959	30	924	3.28	3369

# Experimental results

## *BIRCH clustering*



<i>DS</i>	<i>Time</i>	<i>D</i>	<i># Scan</i>	<i>DS</i>	<i>Time</i>	<i>D</i>	<i># Scan</i>
1	11.5	1.87	2	10	13.6	1.87	2
2	10.7	1.99	2	20	12.1	1.99	2
3	11.4	3.95	2	30	12.2	3.99	2




# Experimental results

---

- Page size
  - *When using Phase 4  $P$  can vary from 256 to 4096 with out much effect on the final results*
- Memory vs. Time
  - *Results generated with low memory can be compensated for by multiple iterations of Phase 4*
- Scalability

# Conclusions and practical use

- Pixel classification in images
- From top to bottom:
  - *BIRCH classification*
  - *Visible wavelength band*
  - *Near-infrared band*





## Conclusions and practical use

---

- Image compression using vector quantization
- Generate codebook for frequently occurring patterns
- BIRCH performs faster than CLARANS or LBG, while getting better compression and nearly as good quality



## Conclusions and practical use

---

- BIRCH works with very large data sets
- Explicitly bounded by computational resources
  - *Runs with specified amount of memory ( $P$ )*
- Superior to CLARANS and KMEANS
  - *Quality, speed, stability and scalability*





# CF tree

---

- CF tree is height balanced tree CF (B,T)
  - Branching factor (B)
  - Threshold (T)
  - Node (leaf node and non-leaf node) represents a cluster made up of all the subclusters represented by its entries
    - Non-leaf node contains at most B entries of the form [CF<sub>i</sub>, child<sub>i</sub>]
    - Leaf node contains at most L entries(objects), each entry represents a sub-cluster, which absorbs many data points with diameter (or radius) under the specified threshold value.

# Insertion into a CF tree

- To insert an object into a CF tree:
  1. Identifying the appropriate leaf
  2. Modifying the leaf
    1. Absorb by a leaf entry
    2. Create a new leaf entry
      1. No leaf splitting
      2. Leaf splitting
  3. Modifying the path to the leaf
    1. No leaf/node split
    2. Leaf/node split
  4. Merging refinement
    1. No resplitting
    2. resplitting